

Automating Design of FPGA Switch-Blocks



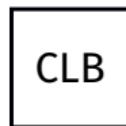
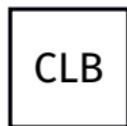
Stefan Nikolić

University of Pennsylvania, 16.02.2023

École Polytechnique Fédérale de Lausanne

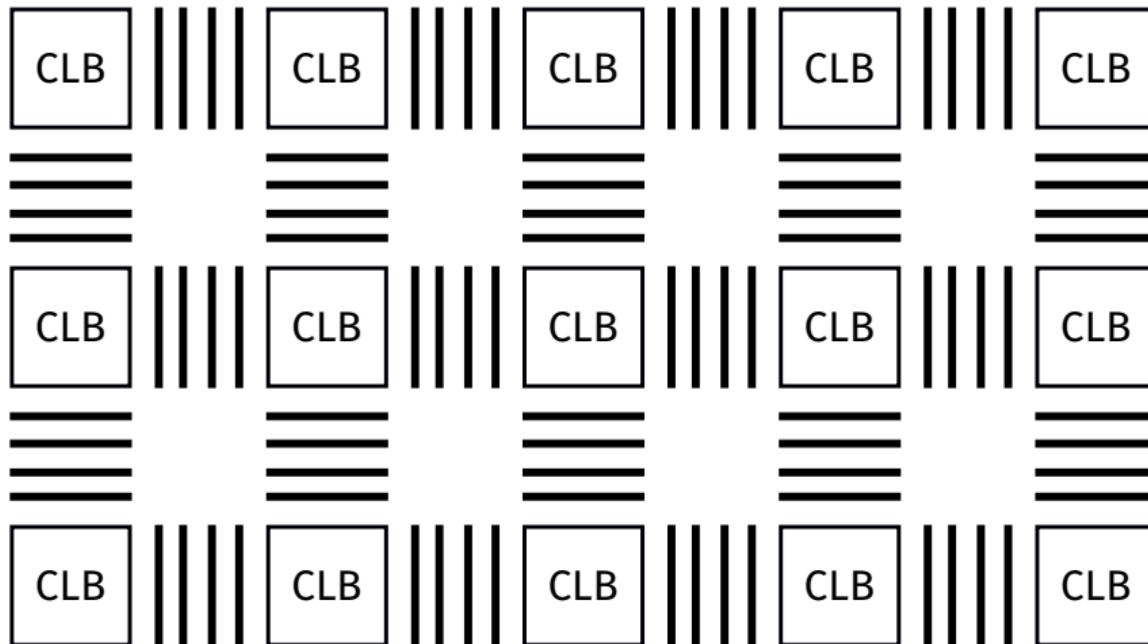
What is the problem?

A very quick review of Island-Style FPGA architecture



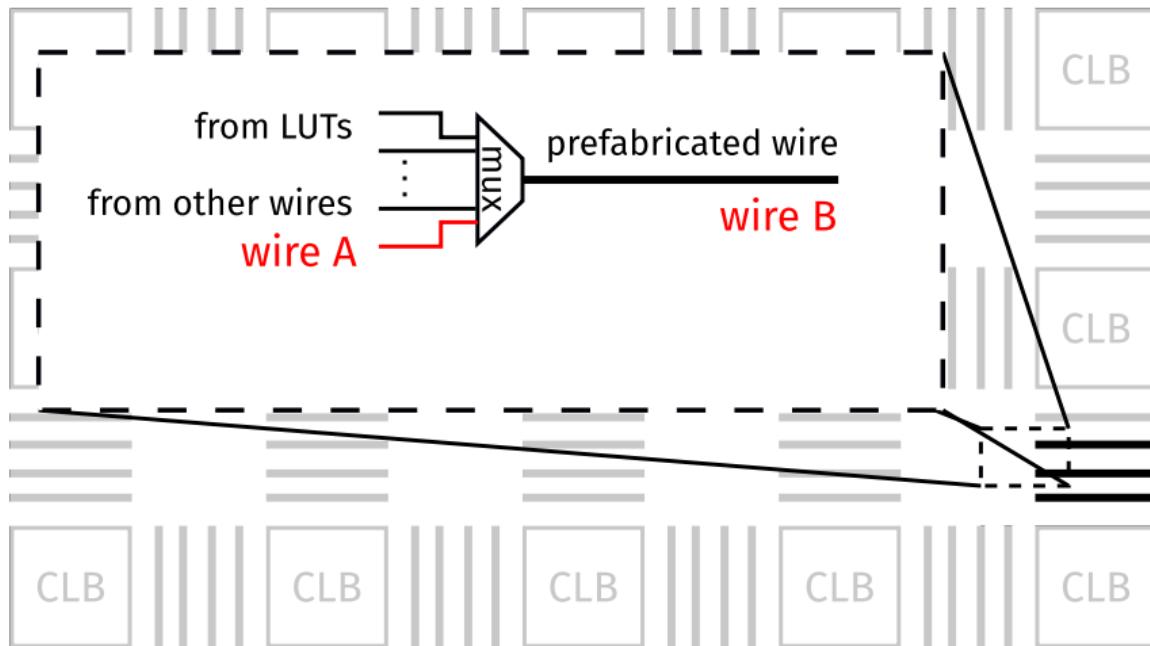
“Islands of LUTs”

A very quick review of Island-Style FPGA architecture



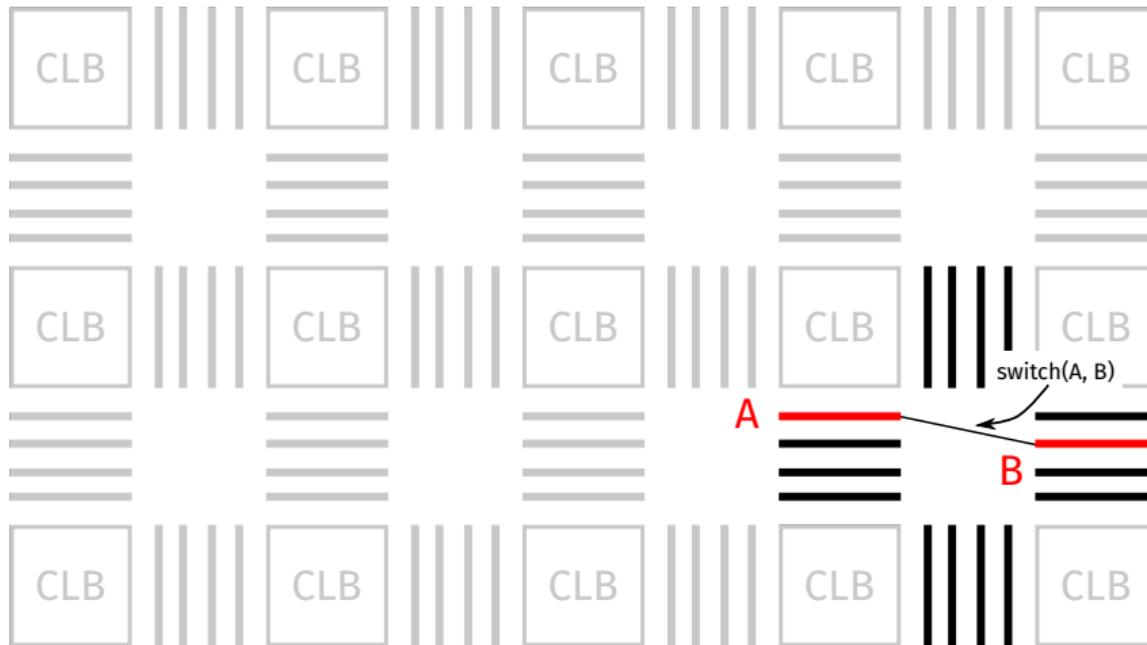
“Surrounded by channels of prefabricated wires”

A very quick review of Island-Style FPGA architecture



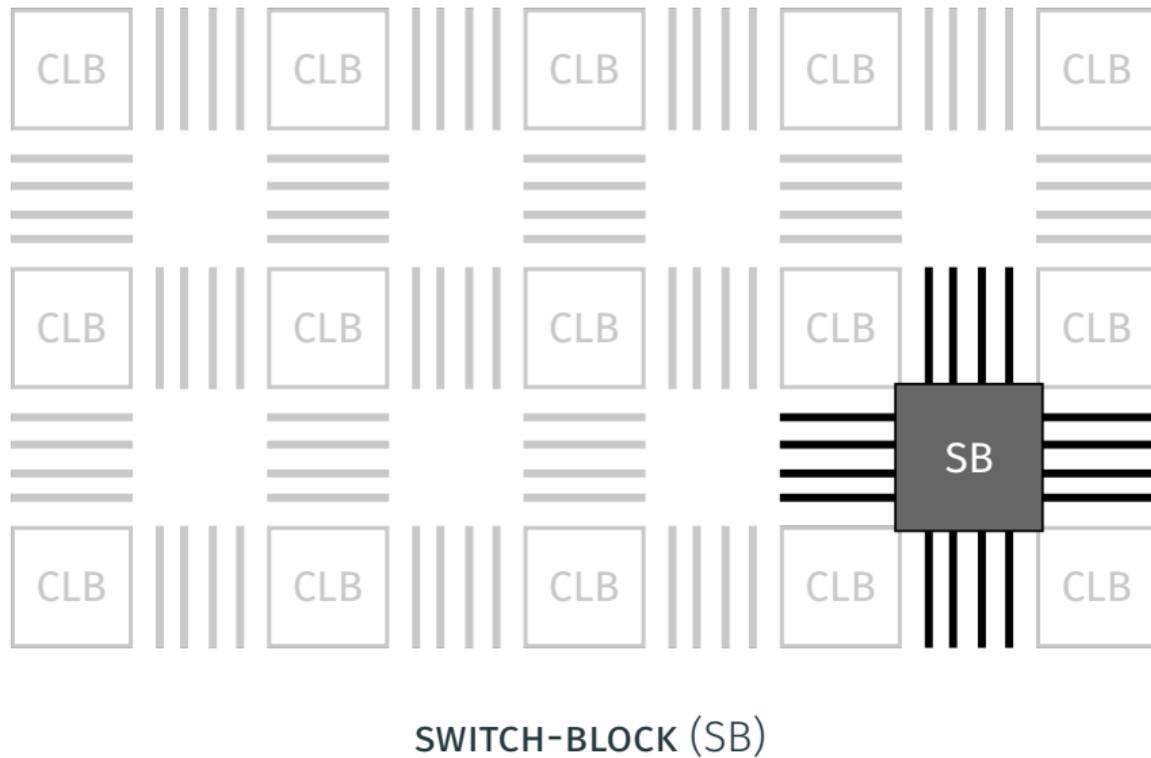
“There is a **SWITCH** between wire A and wire B”

A very quick review of Island-Style FPGA architecture

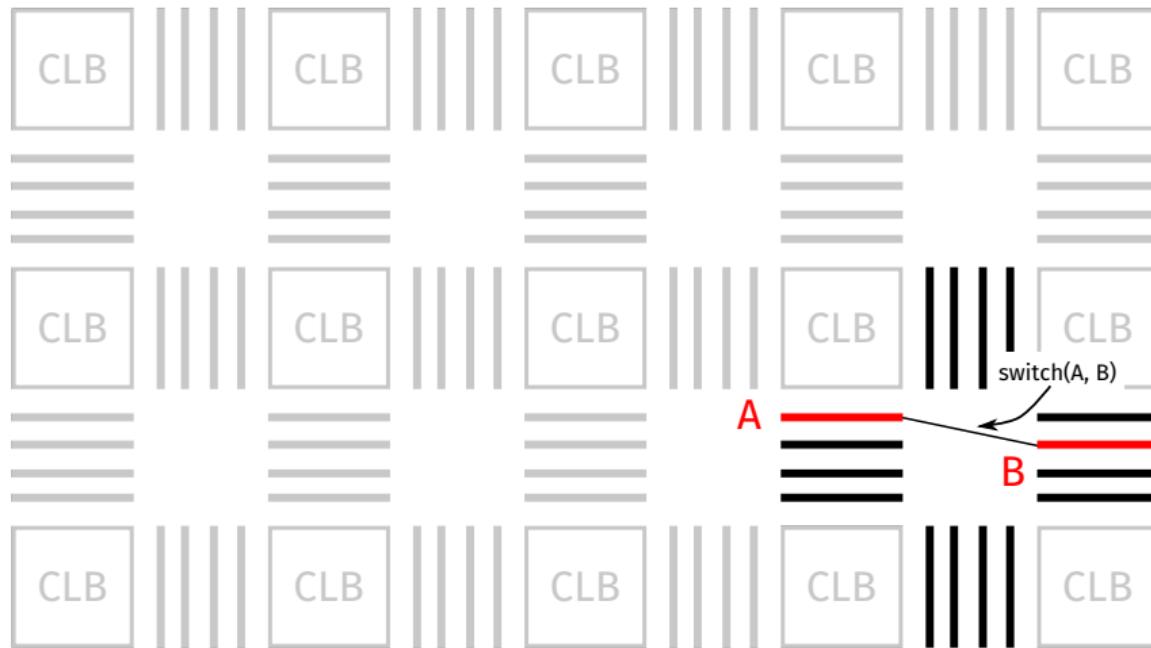


“There is a **SWITCH** between wire A and wire B”

A very quick review of Island-Style FPGA architecture

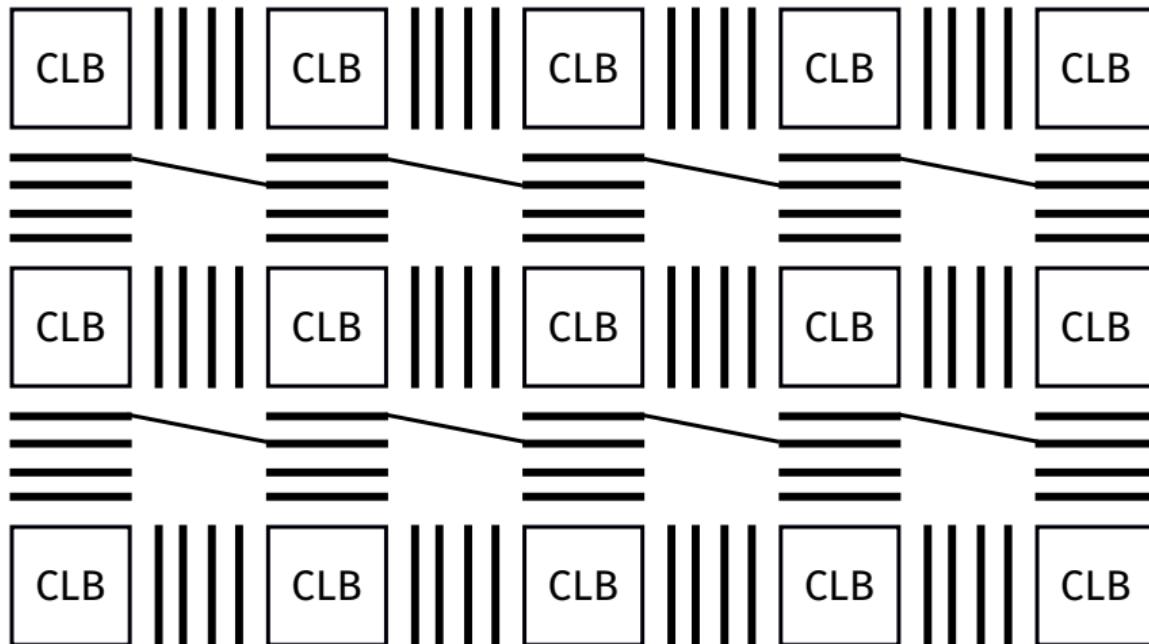


A very quick review of Island-Style FPGA architecture



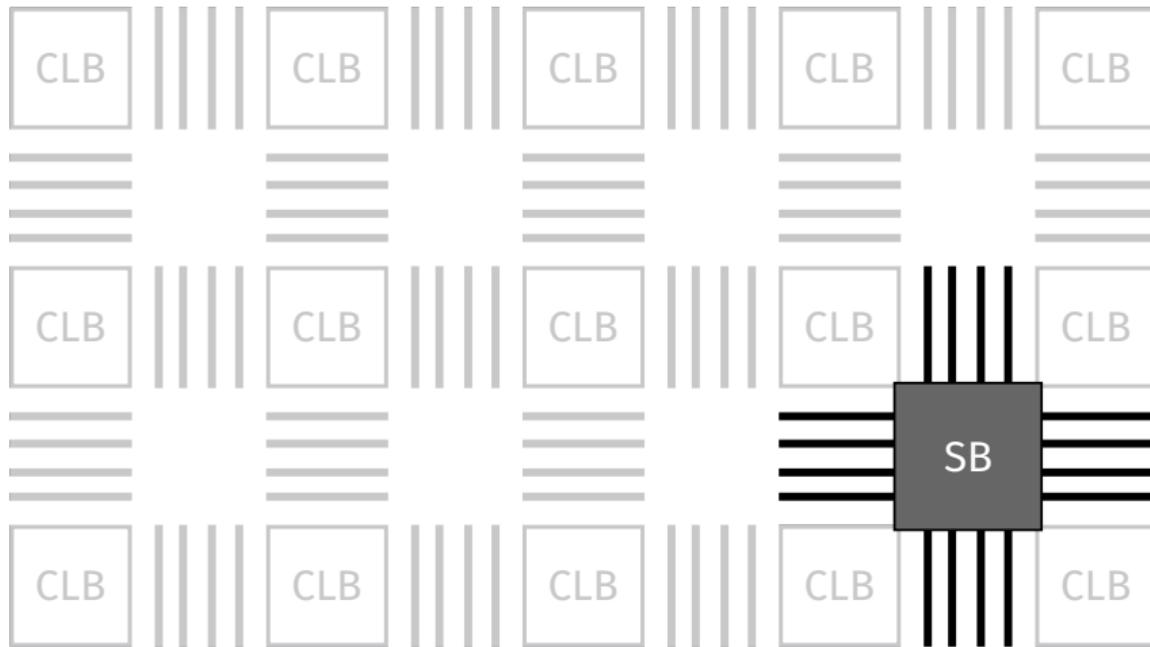
All SBs are identical (e.g., one has switch(A, B) \Rightarrow all have it)

A very quick review of Island-Style FPGA architecture



All SBs are identical (e.g., one has switch(A, B) \Rightarrow all have it)

What is the problem?



Which switches should the SB have?

Isn't switch-pattern design a closed problem?

2019 International Conference on Field-Programmable Technology (ICFPT)

A Study on Switch Block Patterns for Tileable FPGA Routing Architectures

Xifan Tang, Edouard Giacomin, Aurélien Alacchi and Pierre-Emmanuel Gaillardon

University of Utah

Email: xifan.tang@utah.edu

Abstract—Following the rapid growth of *Field Programmable Gate Arrays* (FPGAs) sizes, the regularity of architectures has become a critical feature, leading to the development of million-of-LUT devices. While the routing architecture plays a dominant role in the area, delay and power of modern FPGAs, most of previously published works focus on improving the routability and performance of FPGAs while very few studied tileable (highly-regular) routing architectures. In this paper, we provide a detailed analysis between tileable and popular non-tileable FPGAs considering modern routing architectures. First, we upgrade VPR to generate tileable routing architecture, which can support different switch block patterns for (1) the routing tracks that start/end in a tile and (2) the routing tracks that pass through a tile. Then, we evaluate the performance of mixed switch blocks patterns in the context of a Stratix IV-like FPGA architecture, by considering the most representative patterns, i.e., Subset, Universal and Wilton. Experimental results show that averaged over the MCNC and VTR benchmarks, when compared to the well-optimized non-tileable architectures, the tileable architectures can improve the minimum routable channel width by 13% and area-delay product by 2%. In particular, our results showed that in the context of tileable FPGA, a mix of Universal and Wilton switch block patterns lead to the best trade-off.

Experimental results show that compared to VPR, our RRG generator can reduce the number of unique tiles by 8.8 \times and 5.5 \times for homogeneous and heterogeneous FPGAs respectively, even considering 128 \times 128 array size.

(2) More than tileable FPGA, our RRG generator also supports different switch block patterns for (a) the routing tracks that start/end in a tile and (b) the routing tracks that pass a tile. We evaluate the performance of mixed switch blocks patterns in the context of a Stratix IV-like FPGA architecture, by considering the most representative patterns, i.e., Subset [20], Universal [19], Wilton [21] and Imran [22]. Experimental results show that averaged over the MCNC and VTR benchmarks, when compared to the well-optimized non-tileable architectures, the tileable architectures can improve the minimum routable channel width by 13% and area-delay product by 2%. In particular, our results showed that in the context of tileable FPGA, a mix of Universal and Wilton switch block patterns leads to the best trade-off in area, delay and routability, while Wilton switch block was the best choice in non-tileable FPGAs.

Isn't switch-pattern design a closed problem?

2019 International Conference on Field-Programmable Technology (ICFPT)

A Study on Switch Block Patterns for Tileable FPGA Routing Architectures

Xifan Tang, Edouard Giacomin, Aurélien Alacchi and Pierre-Emmanuel Gaillardon
University of Utah
Email: xifan.tang@utah.edu

Abstract—Following the rapid growth of *Field Programmable Gate Arrays* (FPGAs) sizes, the regularity of architectures has become a critical feature, leading to the development of million-of-LUT devices. While the routing architecture plays a dominant role in the area, delay and power of modern FPGAs, most of previously published works focus on improving the routability and performance of FPGAs while very few studied tileable (highly-regular) routing architectures. In this paper, we provide a detailed analysis between tileable and popular non-tileable FPGAs considering modern routing architectures. First, we upgrade VPR to generate tileable routing architecture, which can support different switch block patterns for (1) the routing tracks that start/end in a tile and (2) the routing tracks that pass through a tile. Then, we evaluate the performance of mixed switch blocks patterns in the context of a Stratix IV-like FPGA architecture, by considering the most representative patterns, i.e., Subset, Universal and Wilton. Experimental results show that averaged over the MCNC and VTR benchmarks, compared to the well-optimized non-tileable

Experimental results show that compared to VPR, our RRG generator can reduce the number of unique tiles by 8.8 \times and 5.5 \times for homogeneous and heterogeneous FPGAs respectively, even considering 128 \times 128 array size.

(2) More than tileable FPGA, our RRG generator can support different switch block patterns for (a) the routing tracks that start/end in a tile and (b) the routing tracks that pass a tile. We evaluate the performance of mixed switch blocks patterns in the context of a Stratix IV-like FPGA architecture, by considering the most representative patterns, i.e., **Subset [20]**, **Universal [19]**, **Wilton [21]** and **Imran [22]**. Experimental results show that averaged over the MCNC and VTR benchmarks, when compared to the well-optimized non-tileable architectures, the tileable routing architectures can improve the routability by 13% and area-delay product by 2%. In particular, our results showed that in the context of tileable FPGA, a mix of Universal and Wilton switch block patterns lead to the best trade-off between area and delay. In the context of non-tileable FPGAs, the Wilton switch block was the best choice in non-tileable FPGAs.

Subset [20]

1997

Universal [19] **Wilton [21]** **Imran [22]**

1996, a mix of Universal and Wilton switch blocks lead to the best trade-off between area, delay and routability. The Wilton switch block was the best choice in non-tileable FPGAs.

Academic assumptions

VTR 7.0: Next Generation Architecture and CAD System for FPGAs

JASON LUU, University of Toronto
JEFFREY GOEDERS, University of British Columbia
MICHAEL WAINBERG, University of Toronto
ANDREW SOMERVILLE, University of New Brunswick
THIEN YU, University of Toronto
KONSTANTIN NASARTSCHUK, University of New Brunswick
MIAD NASR, University of Toronto
SEN WANG, University of New Brunswick
TIM LIU and NOORUDDIN AHMED, University of Toronto
KENNETH B. KENT, University of New Brunswick
JASON ANDERSON, JONATHAN ROSE, and VAUGHN BETZ, University of Toronto

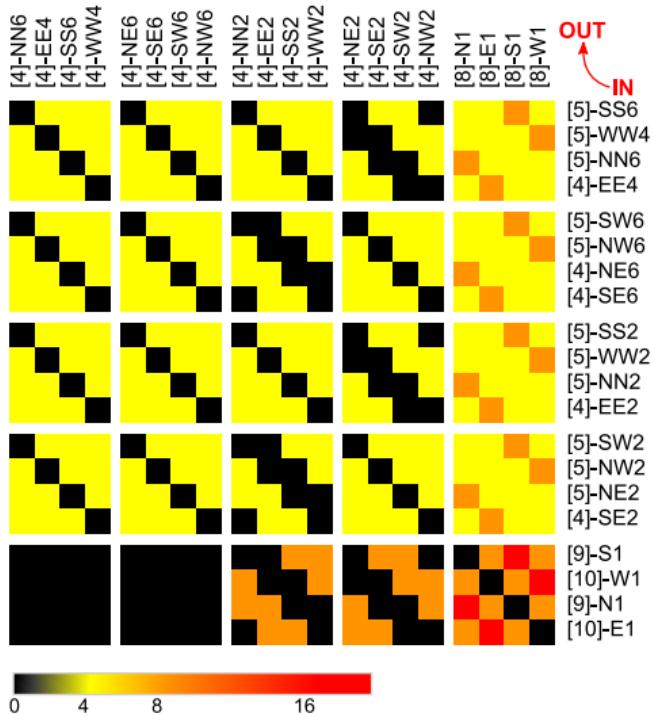
6

3.1. Comprehensive Architecture File

The *Comprehensive Architecture* is the flagship architecture of the VTR release, and we suggest that new users of VTR begin with it. It describes an FPGA architecture with a number of modern features, including fracturable LUTs, carry-chains, fracturable multipliers, and configurable memories. Its architecture features are chosen to be in line with both the recommendations of prior research and current commercial practice in Xilinx's Virtex 7 and Altera's Stratix IV architectures. The routing architecture uses length-4 wire segments, single-driver routing, an F_{Cin} of 0.15, an F_{Cout} of 0.1, and an F_s of 3.

ACM Transactions on Reconfigurable Technology and Systems, Vol. 7, No. 2, Article 6, Publication date: June 2014.

Industrial reality



H1, H2, H4, H12
V1, V2, V6, V12, V18

Fs = 16

Technology driving change

N16: Wu et al.,

"A 16nm FinFET CMOS technology for mobile SoC and computing applications",

IEDM'13

N7: Wu et al.,

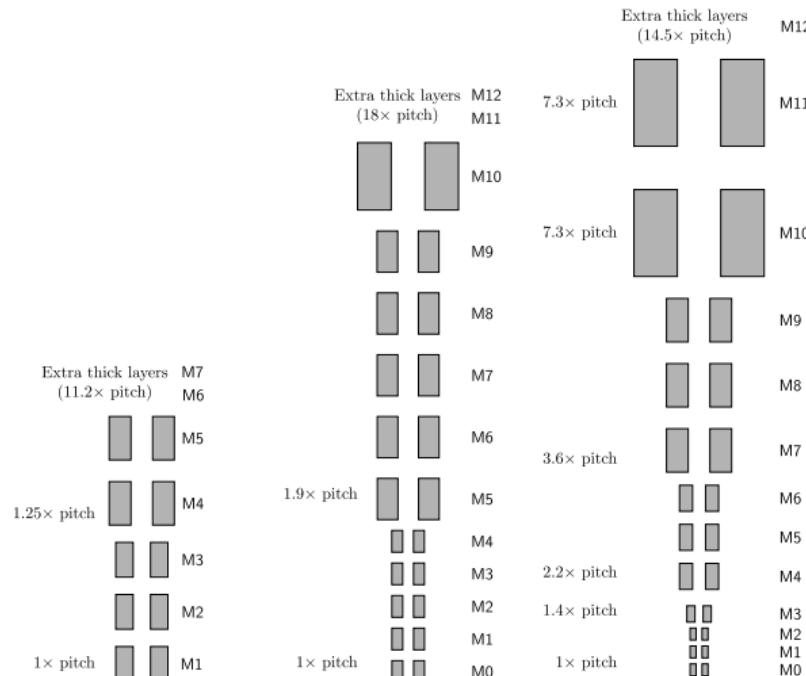
"A 7nm CMOS platform technology featuring 4th generation FinFET transistors with a 0.027 um² high density 6-T SRAM cell for mobile SoC applications",

IEDM'16

N3: Prasad et al.,

"Buried power rails and back-side power grids: Arm® CPU power delivery network design beyond 5nm",

IEDM'19



N16

N7

N3

Technology driving change

N16: Wu et al.,

"A 16nm FinFET CMOS technology for mobile SoC and computing applications",

IEDM'13

N7: Wu et al.,

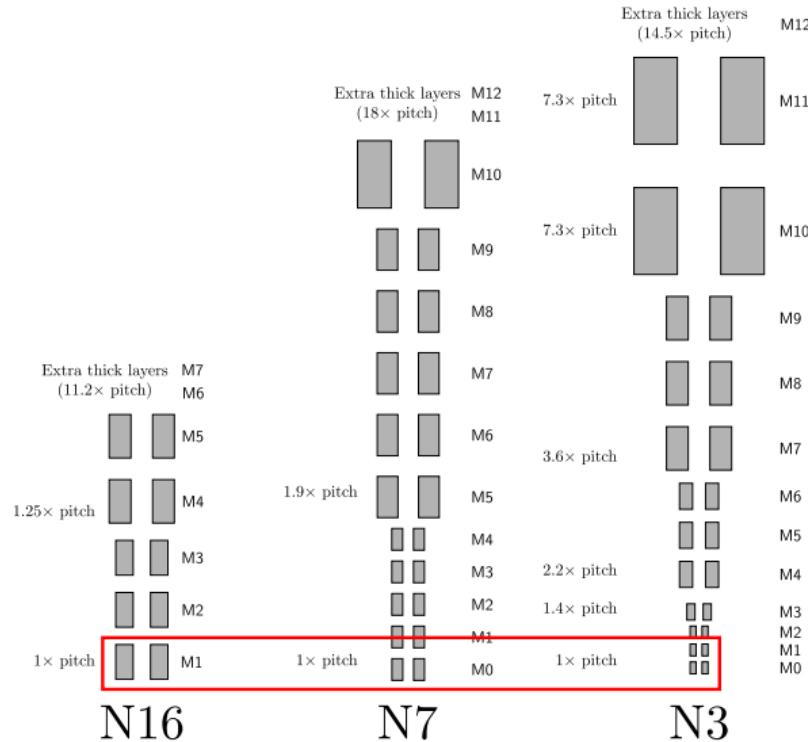
"A 7nm CMOS platform technology featuring 4th generation FinFET transistors with a 0.027 um² high density 6-T SRAM cell for mobile SoC applications",

IEDM'16

N3: Prasad et al.,

"Buried power rails and back-side power grids: Arm® CPU power delivery network design beyond 5nm",

IEDM'19



Technology driving change

Routing Architecture

Intel® Stratix® 10 FPGA

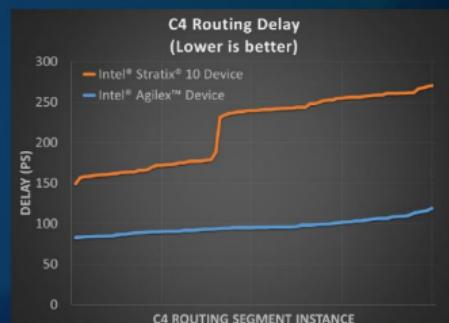
Wide high-fanout MUXes, multi-drop routing segments



Intel® Agilex™ FPGA

Low-fanout, narrow and fast MUXes, single-drop routing segments

Carefully designed routing pattern to maintain and improve routability



[1] Ganusov and Iyer. Agilex Generation of Intel FPGAs. Hot Chips'20

Technology driving change

Routing Architecture

Intel® Stratix® 10 FPGA

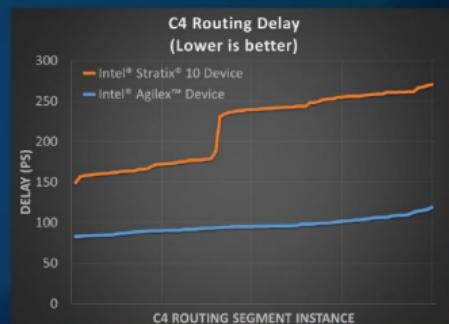
Wide high-fanout MUXes, multi-drop routing segments



Intel® Agilex™ FPGA

Low-fanout, narrow and fast MUXes, single-drop routing segments

Carefully designed routing pattern to maintain and improve routability



“Carefully designed routing pattern to maintain and improve routability”

Two dilemmas

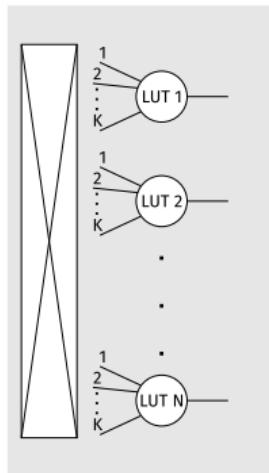
- How to catch up with industry?
- Shouldn't academia set the pace? (like it did 20 years ago)

How do we explore FPGA
architecture?

The Effect of LUT and Cluster Size on Deep-Submicron FPGA Performance and Density

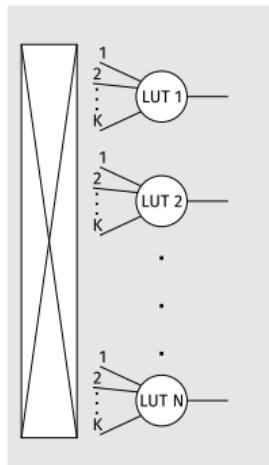
Elias Ahmed and Jonathan Rose

Classical FPGA Architecture Exploration



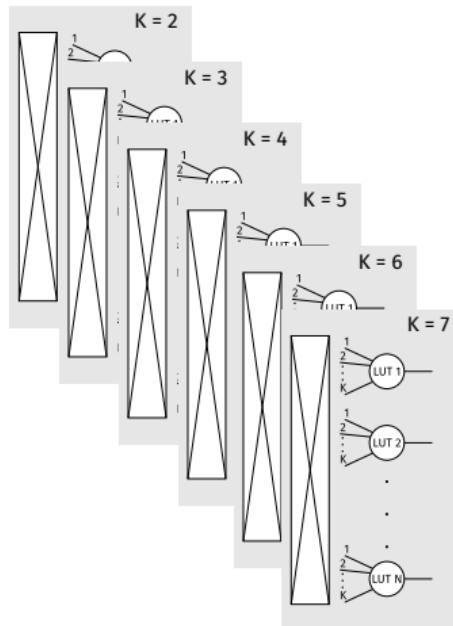
$$K = ?$$

Classical FPGA Architecture Exploration



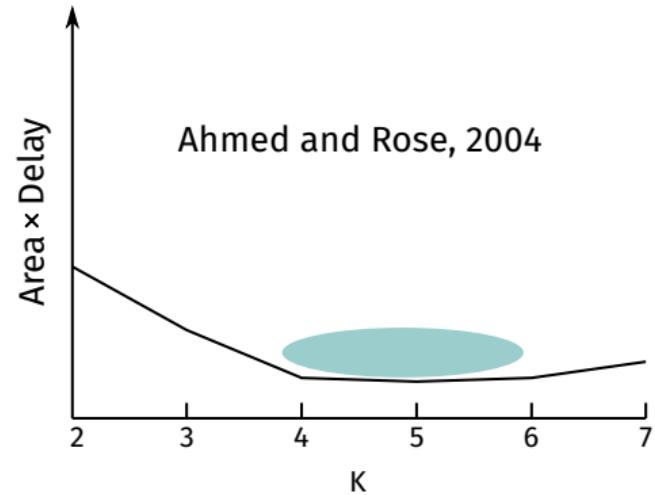
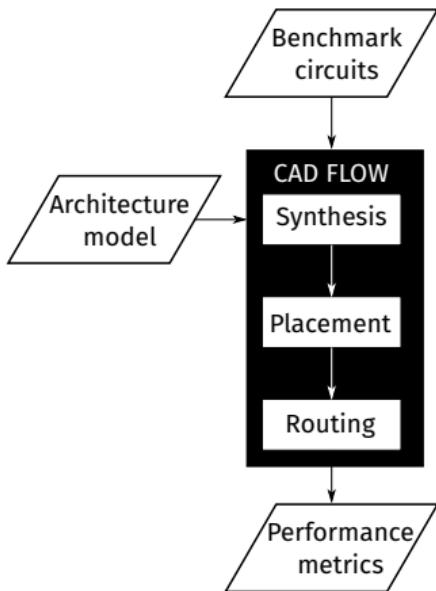
$K = ?$

$$K \in \{2, 3, 4, 5, 6, 7\}$$



Architecture model

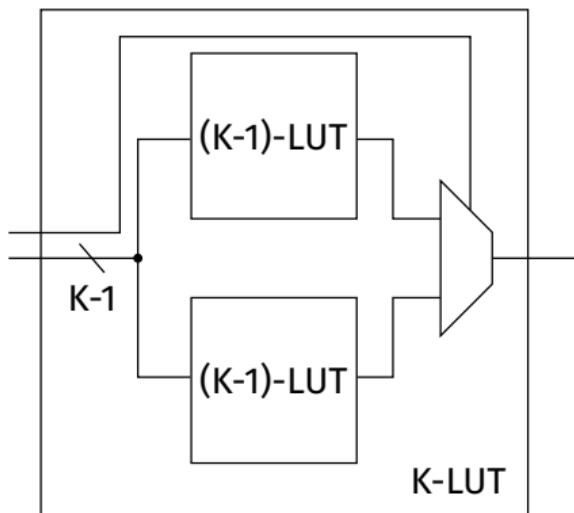
Classical FPGA Architecture Exploration



Ahmed and Rose. The Effect of LUT and Cluster Size on Deep-Submicron FPGA Performance and Density, TVLSI, 2004

Classical FPGA Architecture Exploration

What about $K > 7$?



$$A(K\text{-LUT}) > 2 \times A((K - 1)\text{-LUT})$$

Classical FPGA Architecture Exploration

Similar bounds exist for

- cluster size (crossbar area $\Theta(N^2)$)
- channel wires ($t_{L_{max}} < t_{L'} + t_{L''}, \forall L', L'' < L_{max}$)

How did academic research lead the way?

Academic research has been able to quickly explain commercial FPGA architecture design choices and even drive innovation (e.g., Stratix) through **DESIGN SPACE EXPLORATION**

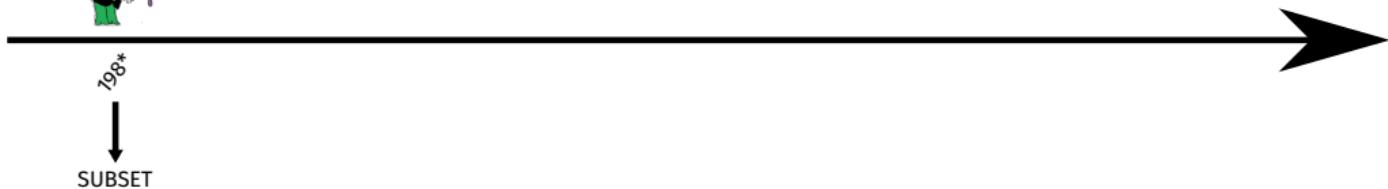
Classical FPGA Architecture Exploration

What about switch-blocks?

Classical FPGA Architecture Exploration

What about switch-blocks?

A very smart engineer at Xilinx
gets an idea



Classical FPGA Architecture Exploration

What about switch-blocks?

A very smart engineer at Xilinx
gets an idea



198*

SUBSET

Some very smart people in
Texas and Hong Kong
get an idea



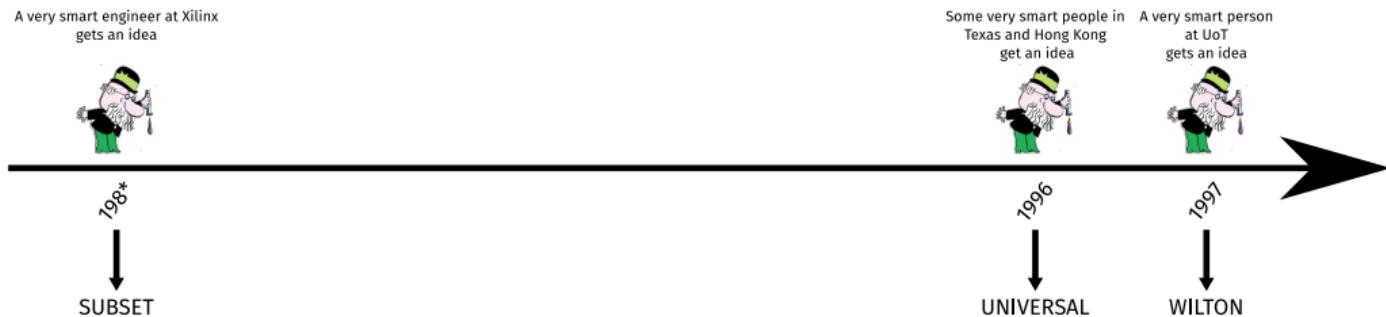
1996

UNIVERSAL



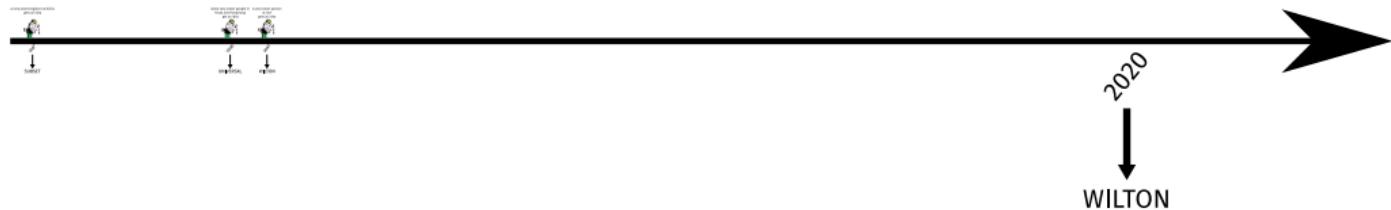
Classical FPGA Architecture Exploration

What about switch-blocks?

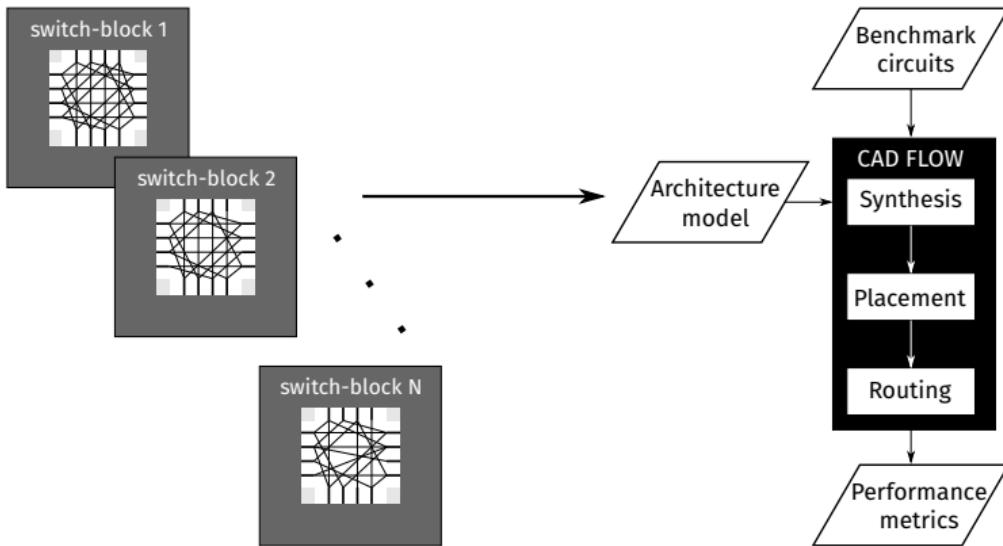


Classical FPGA Architecture Exploration

What about switch-blocks? (apologies for a bit of an exaggeration)



Can't we automate switch-block exploration too?



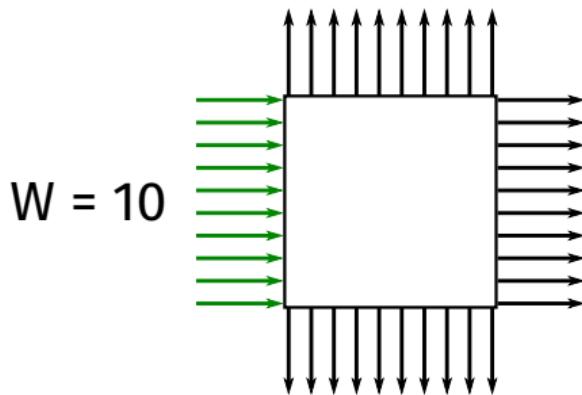
Sure!

Can't we automate switch-block exploration too?

How big is N?

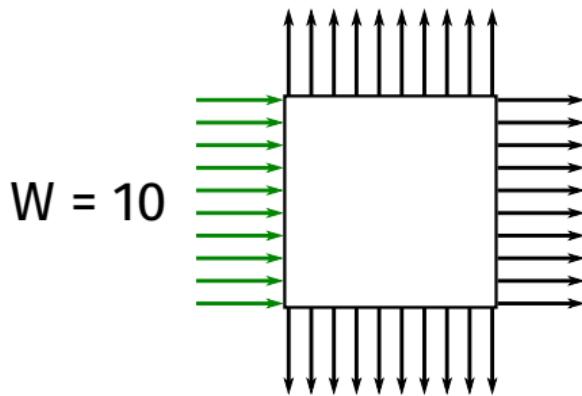
Can't we automate switch-block exploration too?

How big is N?



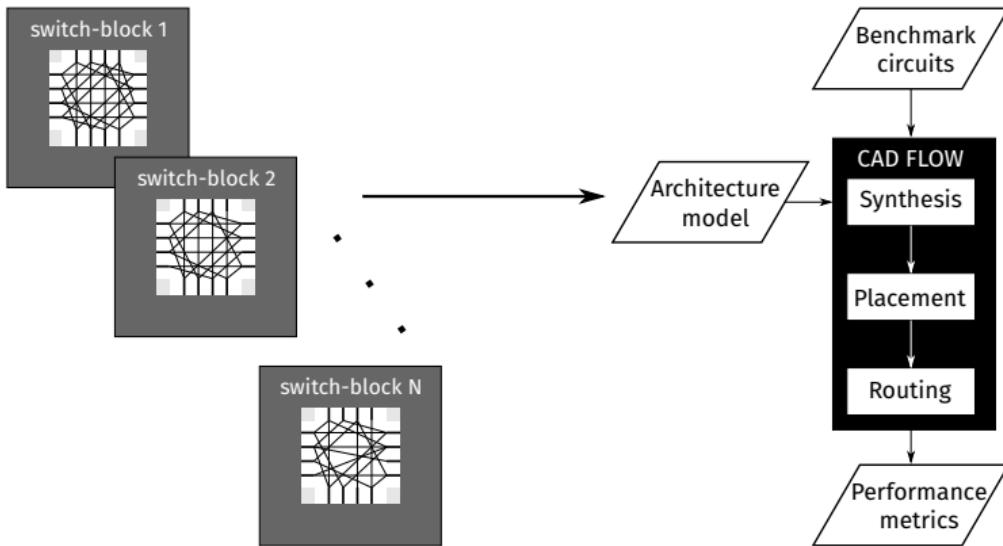
Can't we automate switch-block exploration too?

How big is N?



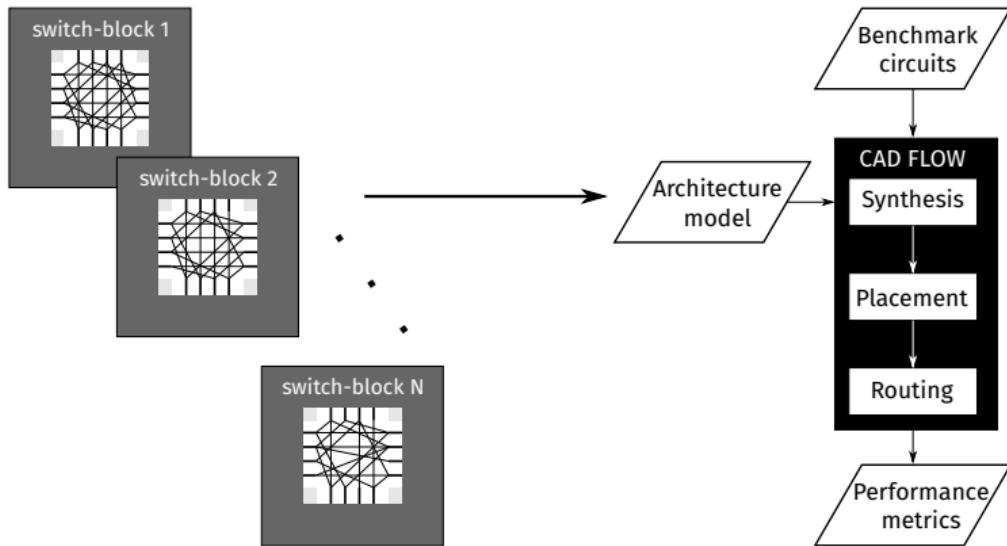
$$2^{10 \times 30}$$

Can't we automate switch-block exploration too?



Sure!

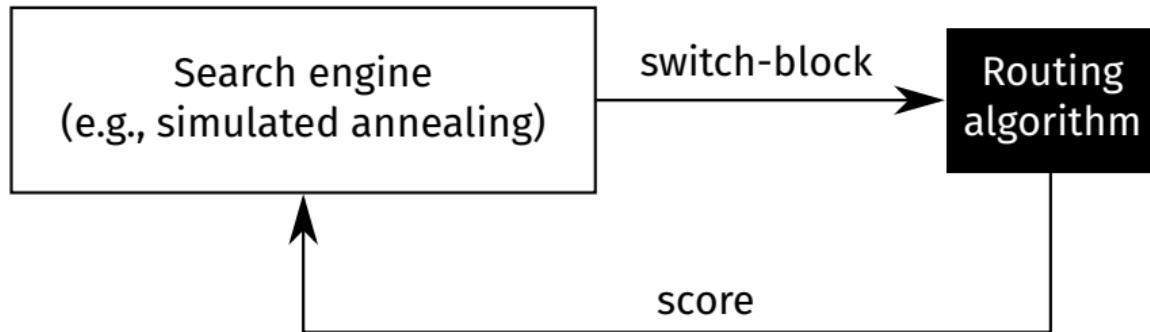
Can't we automate switch-block exploration too?



Surely not like this!

Avalanche Search

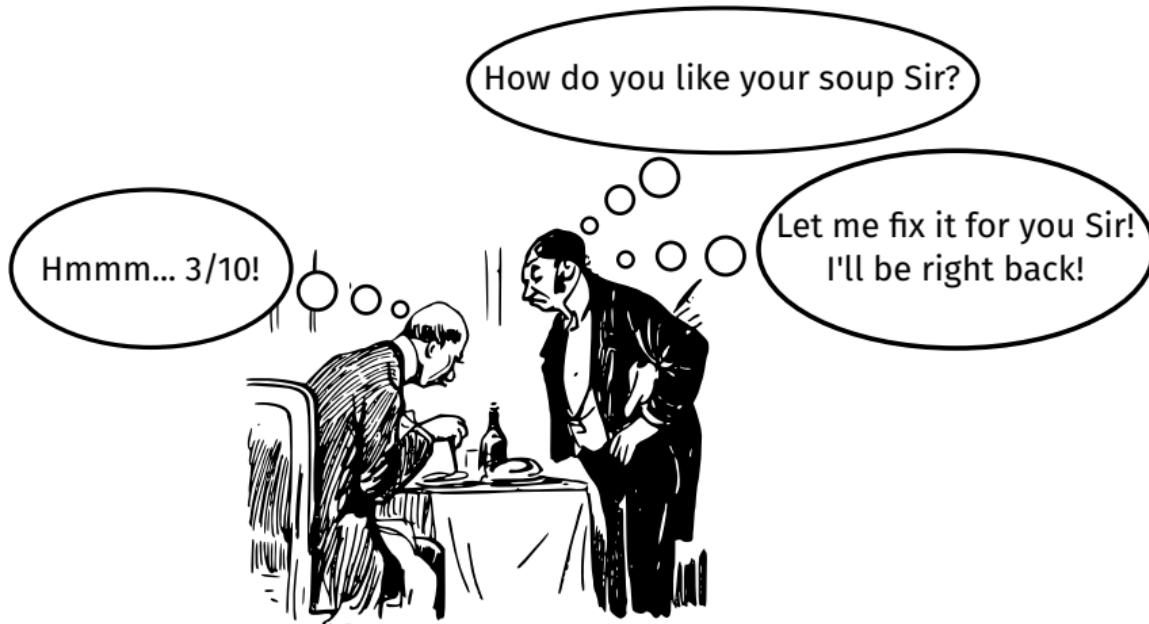
An Intuitive Solution: Iterative Improvement



[1] Lin, Wawrzynek, El Gamal. Exploring FPGA routing architecture stochastically. TCAD'10

Using the router as a black box
to evaluate enumerated solutions
is inefficient

Imagine this restaurant service...



Some things are best left to the consumer



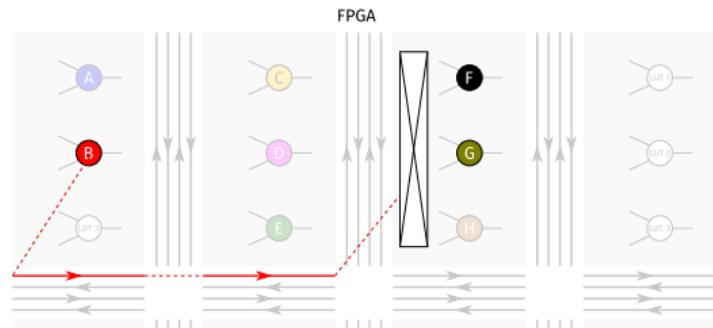
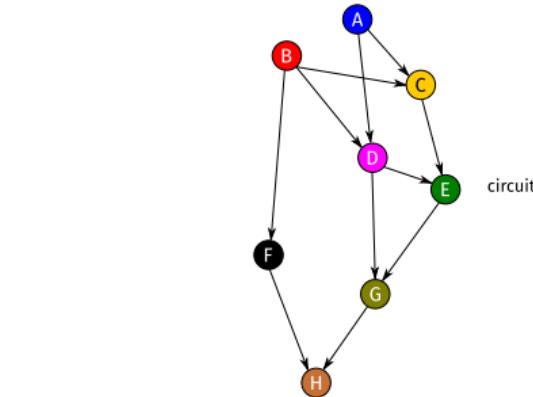
Some things are best left to the consumer



Consumer of switch-block connectivity patterns is the routing algorithm

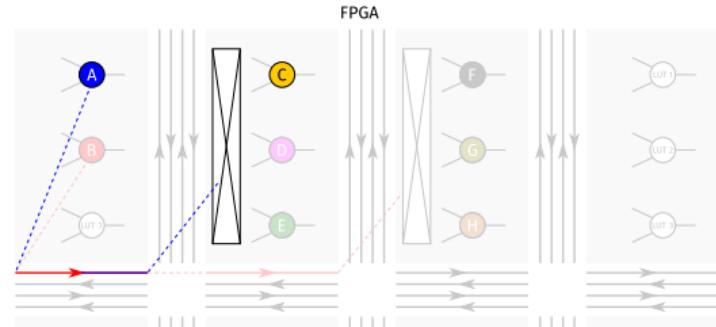
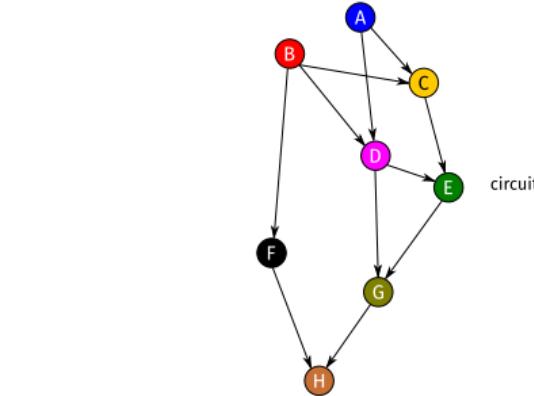
A very brief review of PathFinder

```
1: for  $(u, v) \in E$  do  
2:   path[ $(u, v)$ ] = shortest_path( $u, v$ )  
3: end for
```



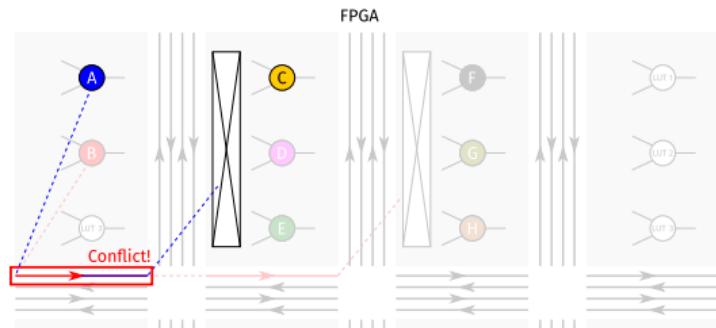
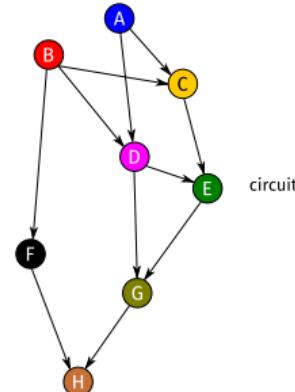
A very brief review of PathFinder

```
1: for  $(u, v) \in E$  do  
2:   path[ $(u, v)$ ] = shortest_path( $u, v$ )  
3: end for
```



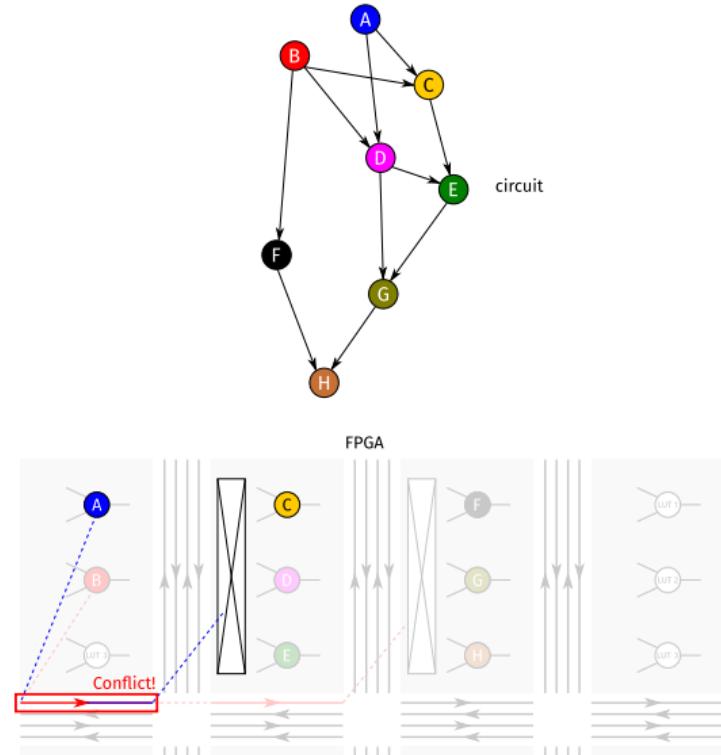
A very brief review of PathFinder

```
1: for  $(u, v) \in E$  do  
2:   path[ $(u, v)$ ] = shortest_path( $u, v$ )  
3: end for
```



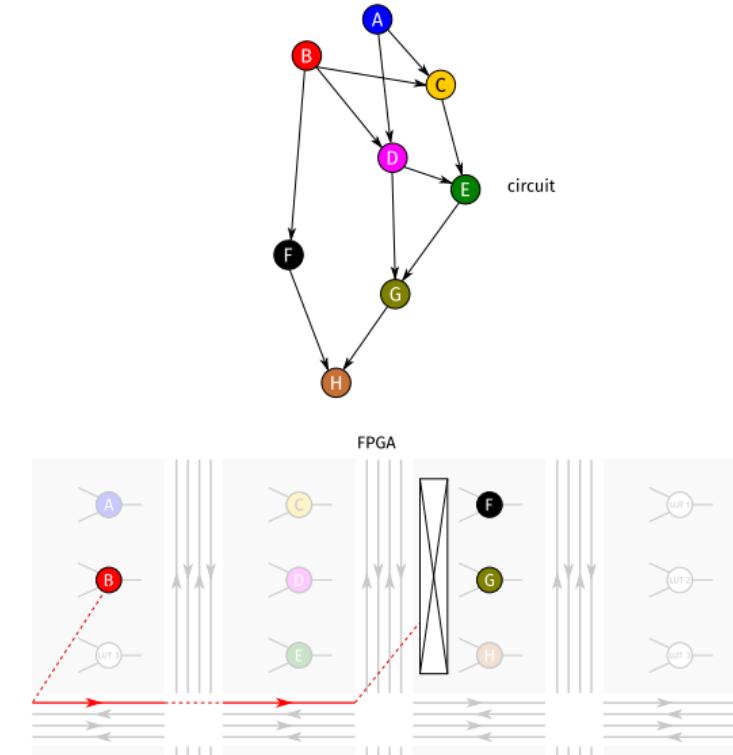
A very brief review of PathFinder

```
1: for  $(u, v) \in E$  do  
2:   path $[(u, v)] = \text{shortest\_path}(u, v)$   
3:   cost $(w) += \alpha, \forall w \in \text{path}[(u, v)]$   
4: end for
```



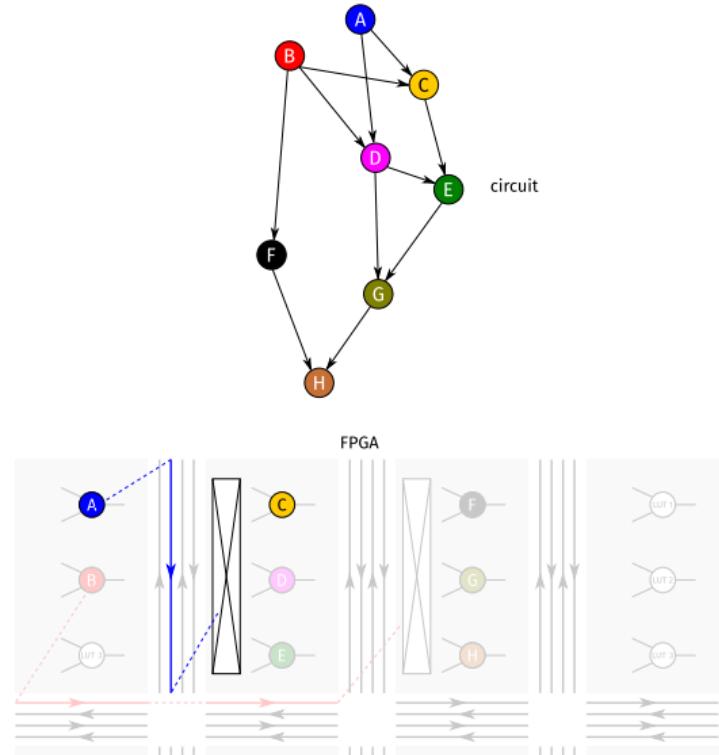
A very brief review of PathFinder

```
1: for  $(u, v) \in E$  do  
2:   path $[(u, v)] = \text{shortest\_path}(u, v)$   
3:   cost $(w) += \alpha, \forall w \in \text{path}[(u, v)]$   
4: end for
```



A very brief review of PathFinder

```
1: for  $(u, v) \in E$  do  
2:   path $[(u, v)] = \text{shortest\_path}(u, v)$   
3:   cost $(w) += \alpha, \forall w \in \text{path}[(u, v)]$   
4: end for
```



Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal

```

1: function CONGESTION_COST( $u, s$ )           ▷ computes the congestion cost of node  $u$  when routing signal  $s$ 
2:   if  $u \in RT(s)$  then return 0           ▷ if  $u$  is already used by one connection of  $s$ , it can freely be used by another
3:   return  $b(u) \times (1 + p_{fac} \times O(u)) \times (1 + h_{fac} \times C_h(u))$            ▷ otherwise, account for congestion
4: for  $u \in V$  do
5:    $O(u) = 0; C_h(u) = 0$                    ▷ set occupancy and historical congestion of all nodes to 0
6:   if  $(\exists v \in V) ((u, v) \in E_c)$  then
7:      $RT(u) = \{u\}$                       ▷ initialize the routing tree of each signal
8:    $i = 0; p_{fac} = p_{fac}^{init}$ 
9:   do
10:    if  $i \geq \text{max\_iter}$  then return UNROUTABLE ▷ no congestion-free routing was found in max_iter iterations
11:    for  $s \in \{u \in V : (\exists v \in V) ((u, v) \in E_c)\}$  do           ▷ all signals are ripped up and rerouted in each iteration;
12:      for  $u \in RT(s)$  do           ▷ modern incremental routers deviate from this [16]
13:         $O(u) = O(u) - 1$            ▷ reduce the occupancy of all nodes used by the signal  $s$  that is ripped up
14:         $RT(s) = \{s\}; O(s) = O(s) + 1$            ▷ rip up the signal
15:        for  $t \in V : (s, t) \in E_c$  do
16:           $P = \text{SHORTEST\_PATH}(s, t, \forall u \in V : \text{cong}(u) = \text{CONGESTION\_COST}(u, s))$  ▷ (re)route the connection  $s \rightarrow t$ 
17:          for  $u \in P$  do
18:            if  $\neg(u \in RT(s))$  then
19:               $O(u) = O(u) + 1$            ▷ increase the occupancy of all nodes not already used by the signal  $s$ 
20:             $RT(s) = RT(s) \cup P$            ▷ add the connection route to the routing tree of  $s$ 
21:          for  $u \in V$  do
22:             $C_h(u) = C_h(u) + \max(0, O(u) - 1)$            ▷ update historical congestion
23:           $i = i + 1$ 
24:         $p_{fac} = p_{fac} \times p_{fac}^{mult}$            ▷ increase the penalty of using occupied nodes;  $p_{fac}^{mult} > 1$  (1.3 is default in VPR [16])
25:      while  $\exists u \in V : O(u) > 1$            ▷ finish if there is no congestion
26:    return  $\forall RT$            ▷ return all routing trees

```

Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal

```

1: function CONGESTION_COST( $u, s$ )           ▷ computes the congestion cost of node  $u$  when routing signal  $s$ 
2:   if  $u \in RT(s)$  then return 0           ▷ if  $u$  is already used by one connection of  $s$ , it can freely be used by another
3:   return  $b(u) \times (1 + p_{fac} \times O(u)) \times (1 + h_{fac} \times C_h(u))$            ▷ otherwise, account for congestion

4: for  $u \in V$  do
5:    $O(u) = 0; C_h(u) = 0$                    ▷ set occupancy and historical congestion of all nodes to 0
6:   if  $(\exists v \in V) ((u, v) \in E_c)$  then
7:      $RT(u) = \{u\}$                       ▷ initialize the routing tree of each signal
8:    $i = 0; p_{fac} = p_{fac}^{init}$ 
9:   do
10:    if  $i \geq \text{max\_iter}$  then return UNROUTABLE  ▷ no congestion-free routing was found in max_iter iterations
11:    for  $s \in \{u \in V : (\exists v \in V) ((u, v) \in E_c)\}$  do          ▷ all signals are ripped up and rerouted in each iteration;
12:      for  $u \in RT(s)$  do                                ▷ modern incremental routers deviate from this [16]
13:         $O(u) = O(u) - 1$                          ▷ reduce the occupancy of all nodes used by the signal  $s$  that is ripped up
14:       $RT(s) = \{s\}; O(s) = O(s) + 1$              ▷ rip up the signal
15:      for  $t \in V : (s, t) \in E_c$  do
16:         $P = \text{SHORTEST\_PATH}(s, t, \forall u \in V : \text{cong}(u) = \text{CONGESTION\_COST}(u, s))$  ▷ (re)route the connection  $s \rightarrow t$ 
17:        for  $u \in P$  do
18:          if  $\neg(u \in RT(s))$  then
19:             $O(u) = O(u) + 1$                      ▷ increase the occupancy of all nodes not already used by the signal  $s$ 
20:           $RT(s) = RT(s) \cup P$                   ▷ add the connection route to the routing tree of  $s$ 
21:        for  $u \in V$  do
22:           $C_h(u) = C_h(u) + \max(0, O(u) - 1)$     ▷ update historical congestion
23:         $i = i + 1$ 
24:       $p_{fac} = p_{fac} \times p_{fac}^{mult}$        ▷ increase the penalty of using occupied nodes;  $p_{fac}^{mult} > 1$  (1.3 is default in VPR [16])
25:    while  $\exists u \in V : O(u) > 1$            ▷ finish if there is no congestion
26:    return  $\forall RT$                       ▷ return all routing trees

```

Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal s

```

1: function CONGESTION_COST( $u, s$ )           ▷ computes the congestion cost of node u when routing signal s
2:   if  $u \in RT(s)$  then return 0           ▷ if u is already used by one connection of s, it can freely be used by another
3:   return  $b(u) \times (1 + p_{fac} \times O(u)) \times (1 + h_{fac} \times C_h(u))$            ▷ otherwise, account for congestion

4: for  $u \in V$  do
5:    $O(u) = 0; C_h(u) = 0$                    ▷ set occupancy and historical congestion of all nodes to 0
6:   if  $(\exists v \in V) ((u, v) \in E_c)$  then
7:      $RT(u) = \{u\}$                       ▷ initialize the routing tree of each signal
8:    $i = 0; p_{fac} = p_{fac}^{init}$ 
9:   do
10:    if  $i \geq \text{max\_iter}$  then return UNROUTABLE  ▷ no congestion-free routing was found in max_iter iterations
11:    for  $s \in \{u \in V : (\exists v \in V) ((u, v) \in E_c)\}$  do          ▷ all signals are ripped up and rerouted in each iteration;
12:      for  $u \in RT(s)$  do          ▷ modern incremental routers deviate from this [16]
13:         $O(u) = O(u) - 1$           ▷ reduce the occupancy of all nodes used by the signal s that is ripped up
14:         $RT(s) = \{s\}; O(s) = O(s) + 1$           ▷ rip up the signal
15:        for  $t \in V : (s, t) \in E_c$  do
16:           $P = \text{SHORTEST\_PATH}(s, t, \forall u \in V : cong(u) = \text{CONGESTION\_COST}(u, s))$   ▷ (re)route the connection s → t
17:          for  $u \in P$  do
18:            if  $\neg(u \in RT(s))$  then
19:               $O(u) = O(u) + 1$           ▷ increase the occupancy of all nodes not already used by the signal s
20:             $RT(s) = RT(s) \cup P$           ▷ add the connection route to the routing tree of s
21:           $C_h(u) = C_h(u) + \max(0, O(u) - 1)$           ▷ update historical congestion
22:        for  $u \in V$  do
23:           $C_h(u) = C_h(u) + \max(0, O(u) - 1)$           ▷ update historical congestion
24:         $i = i + 1$ 
25:         $p_{fac} = p_{fac} \times p_{fac}^{mult}$           ▷ increase the penalty of using occupied nodes;  $p_{fac}^{mult} > 1$  (1.3 is default in VPR [16])
26:      while  $\exists u \in V : O(u) > 1$           ▷ finish if there is no congestion
27:      return  $\forall RT$           ▷ return all routing trees

```

Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal s

```

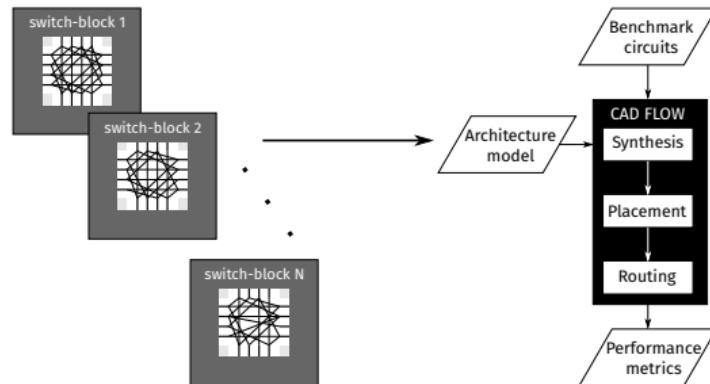
1: function CONGESTION_COST( $u, s$ )           ▷ computes the congestion cost of node  $u$  when routing signal s
2:   if  $u \in RT(s)$  then return 0      ▷ if  $u$  is already used by one connection of  $s$ , it can freely be used by another
3:   return  $b(u) \times (1 + p_{fac} \times O(u)) \times (1 + h_{fac} \times C_h(u))$           ▷ otherwise, account for congestion

4: for  $u \in V$  do
5:    $O(u) = 0; C_h(u) = 0$                   ▷ set occupancy and historical congestion of all nodes to 0
6:   if  $(\exists v \in V) ((u, v) \in E_c)$  then
7:      $RT(u) = \{u\}$                       ▷ initialize the routing tree of each signal
8:    $i = 0; p_{fac} = p_{fac}^{init}$ 
9:   do
10:    if  $i \geq \text{max\_iter}$  then return UNROUTABLE  ▷ no congestion-free routing was found in max_iter iterations
11:    for  $s \in \{u \in V : (\exists v \in V) ((u, v) \in E_c)\}$  do  ▷ all signals are ripped up and rerouted in each iteration;
12:      for  $u \in RT(s)$  do  ▷ modern incremental routers deviate from this [16]
13:         $O(u) = O(u) - 1$                 ▷ reduce the occupancy of all nodes used by the signal  $s$  that is ripped up
14:         $RT(s) = \{s\}; O(s) = O(s) + 1$   ▷ rip up the signal
15:      for  $t \in V : (s, t) \in E_c$  do
16:         $P = \text{SHORTEST\_PATH}(s, t, \forall u \in V : \text{cong}(u) = \text{CONGESTION\_COST}(u, s))$  ▷ (re)route the connection  $s \rightarrow t$ 
17:        for  $u \in P$  do
18:          if  $\neg(u \in RT(s))$  then
19:             $O(u) = O(u) + 1$               ▷ increase the occupancy of all nodes not already used by the signal s
20:           $RT(s) = RT(s) \cup P$           ▷ add the connection route to the routing tree of s
21:        for  $u \in V$  do
22:           $C_h(u) = C_h(u) + \max(0, O(u) - 1)$   ▷ update historical congestion
23:         $i = i + 1$ 
24:       $p_{fac} = p_{fac} \times p_{fac}^{mult}$   ▷ increase the penalty of using occupied nodes;  $p_{fac}^{mult} > 1$  (1.3 is default in VPR [16])
25:      while  $\exists u \in V : O(u) > 1$   ▷ finish if there is no congestion
26:    return VRT  ▷ return all routing trees

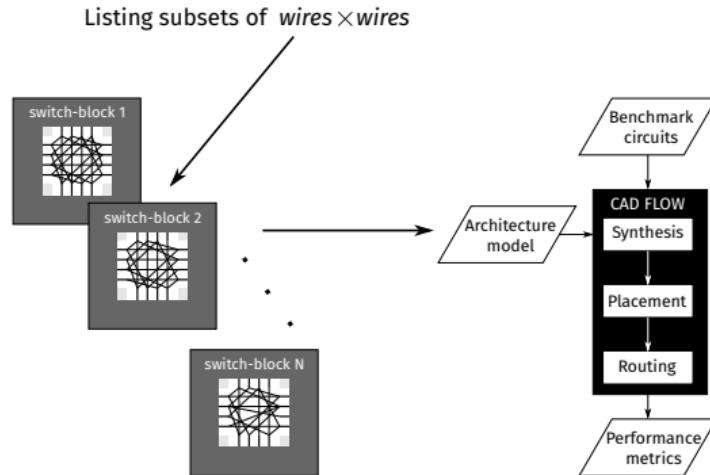
```

Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

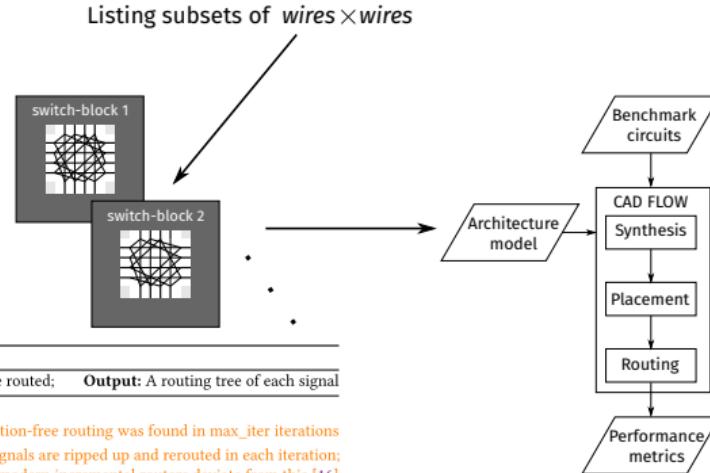
How to avoid listing individual solutions?



How to avoid listing individual solutions?



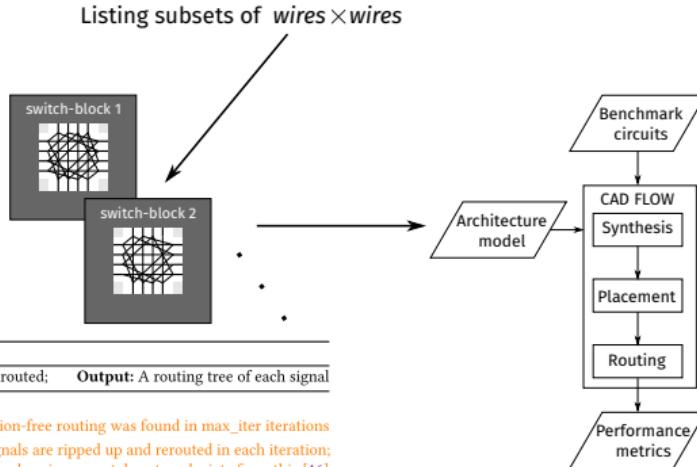
How to avoid listing individual solutions?



Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal

How to avoid listing individual solutions?

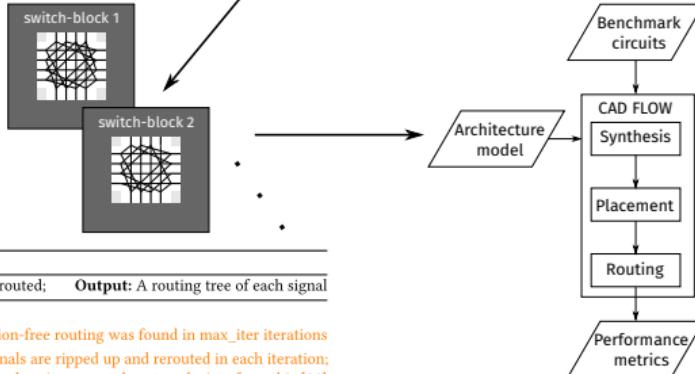


Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal

The more paths use a switch
 $s \in SB \setminus wires \times wires$
 the more useful it is in the architecture

How to avoid listing individual solutions?



Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal

Represent all possible switches

The more paths use a switch

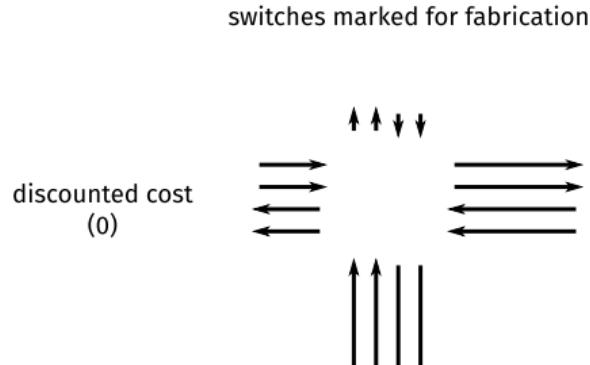
$s \in SB = \text{wires} \times \text{wires}$

the more useful it is in the architecture

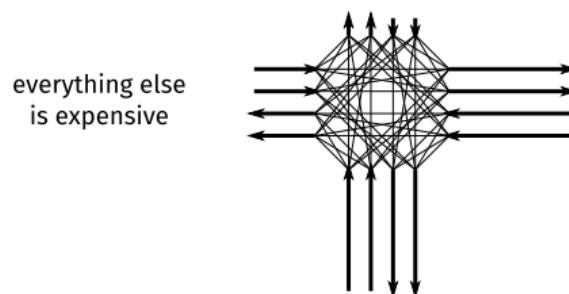
Construct the solution from the most-used ones

1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C

Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]

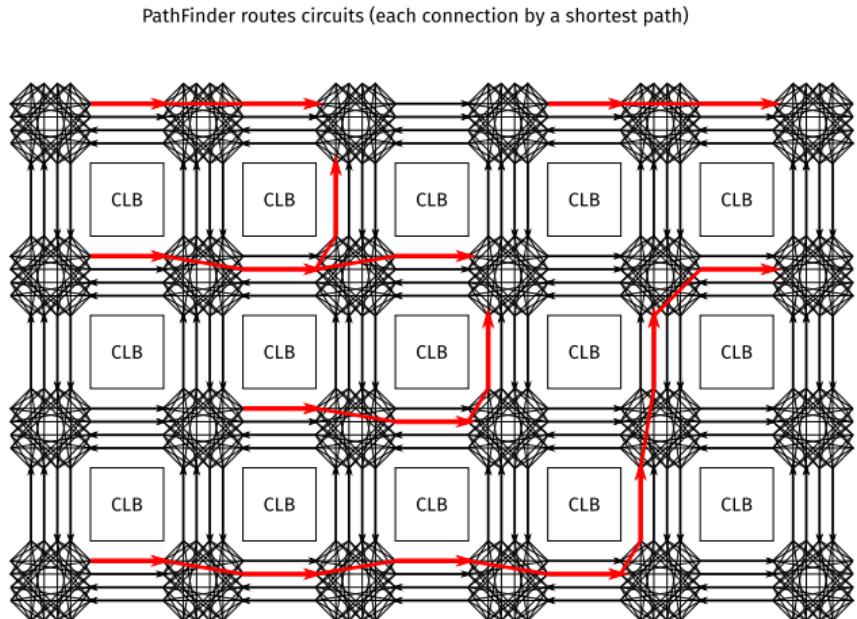
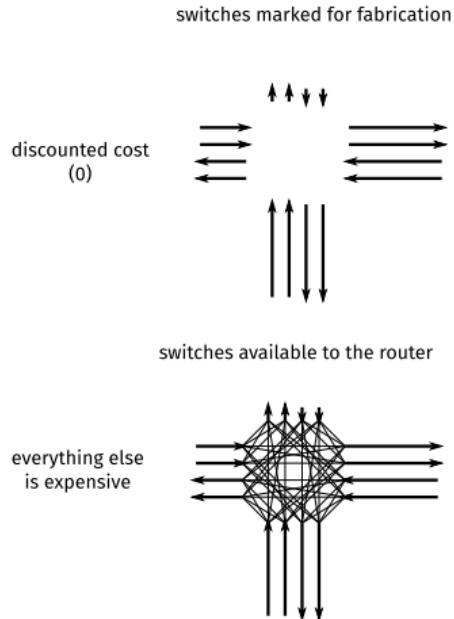


switches available to the router



1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits

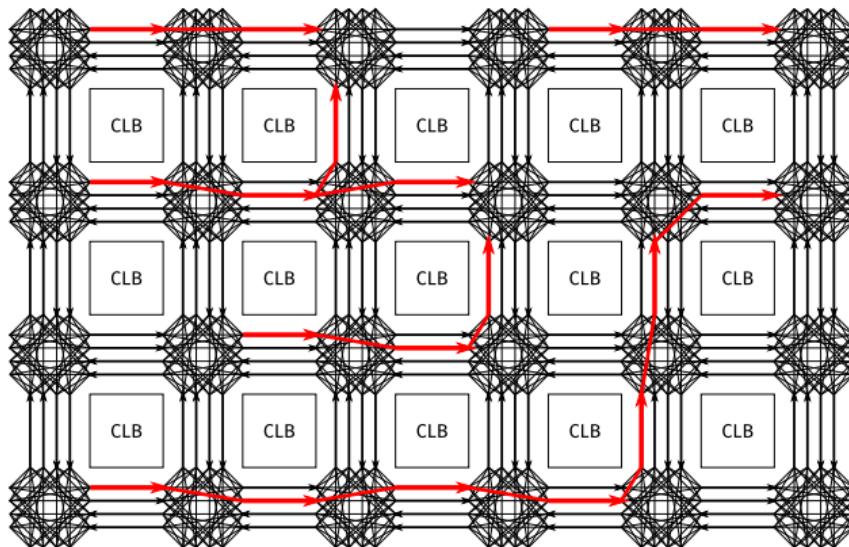
Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]



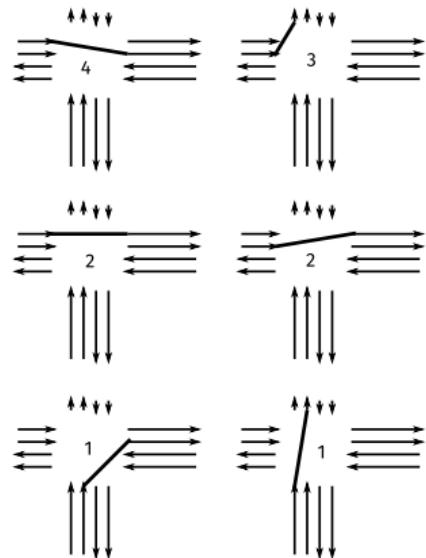
1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done

Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]

PathFinder routes circuits (each connection by a shortest path)



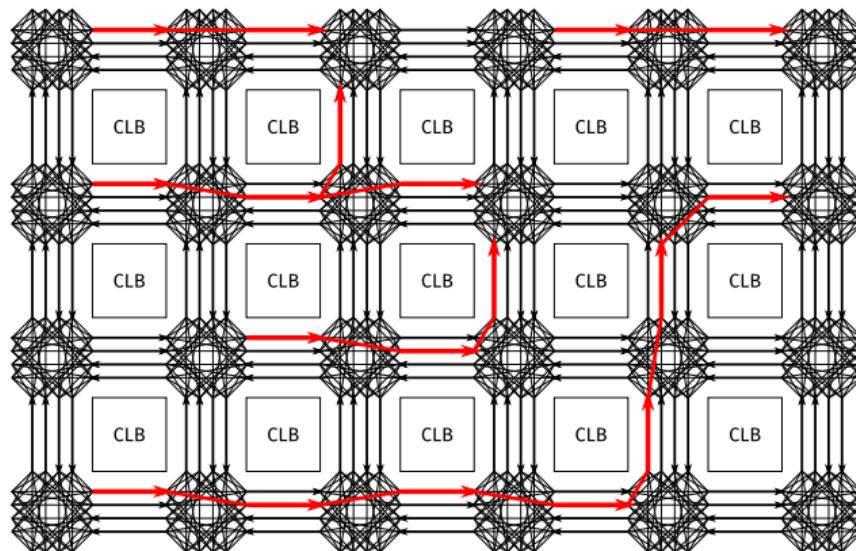
Count in how many switch-blocks each switch was used



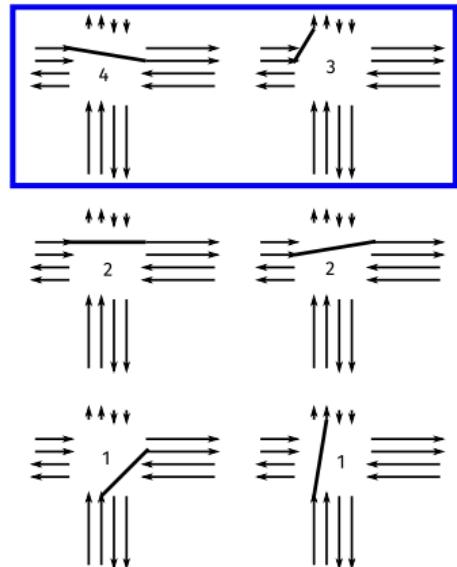
1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done
5. mark n most-used unmarked switches and set their cost to 0

Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]

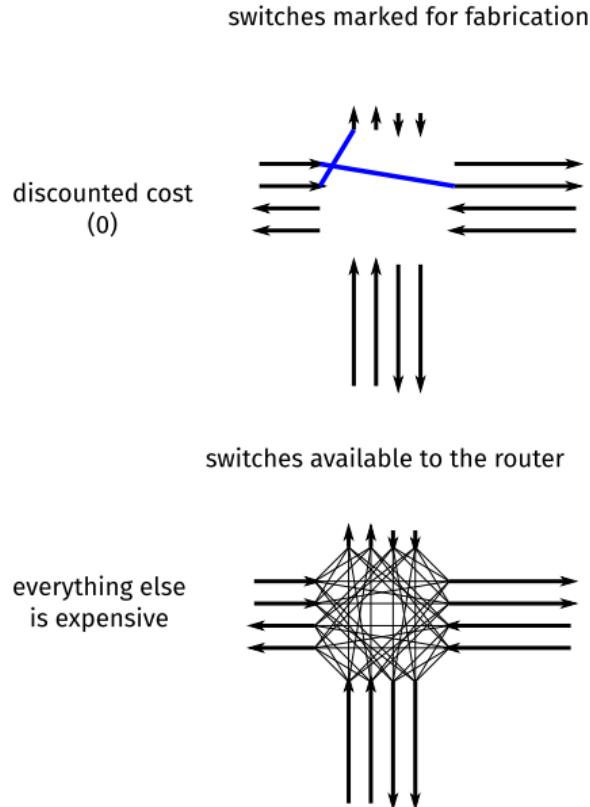
PathFinder routes circuits (each connection by a shortest path)



Give discounts and mark for fabrication

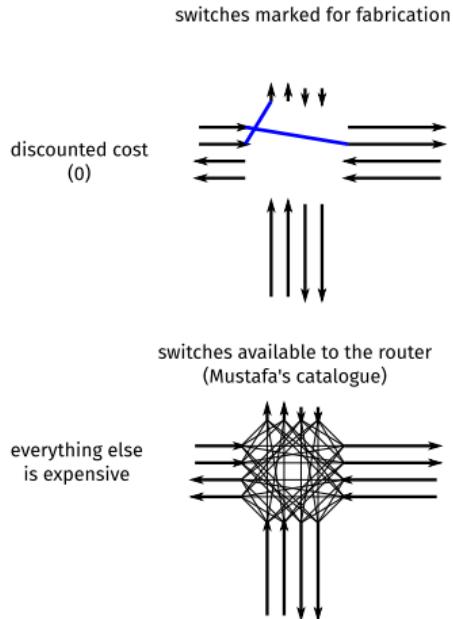


Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]

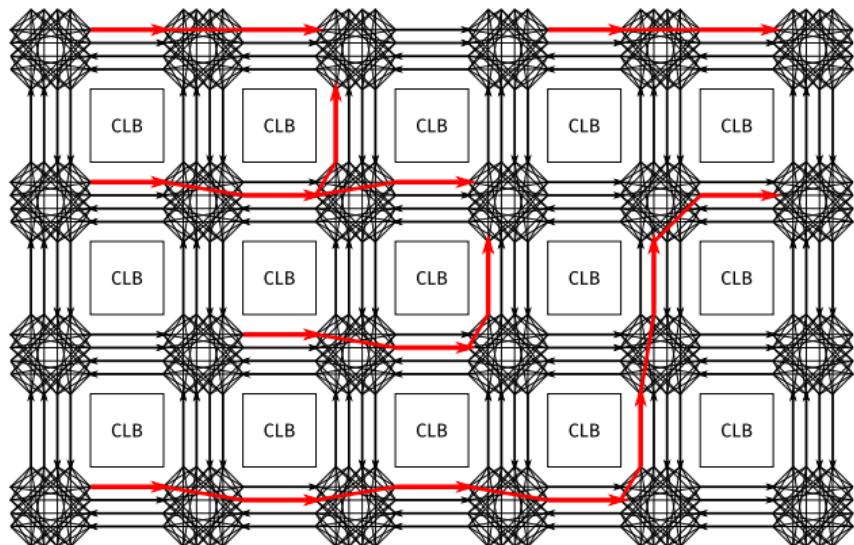


1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done
5. mark n most-used unmarked switches and set their cost to 0
6. goto 3

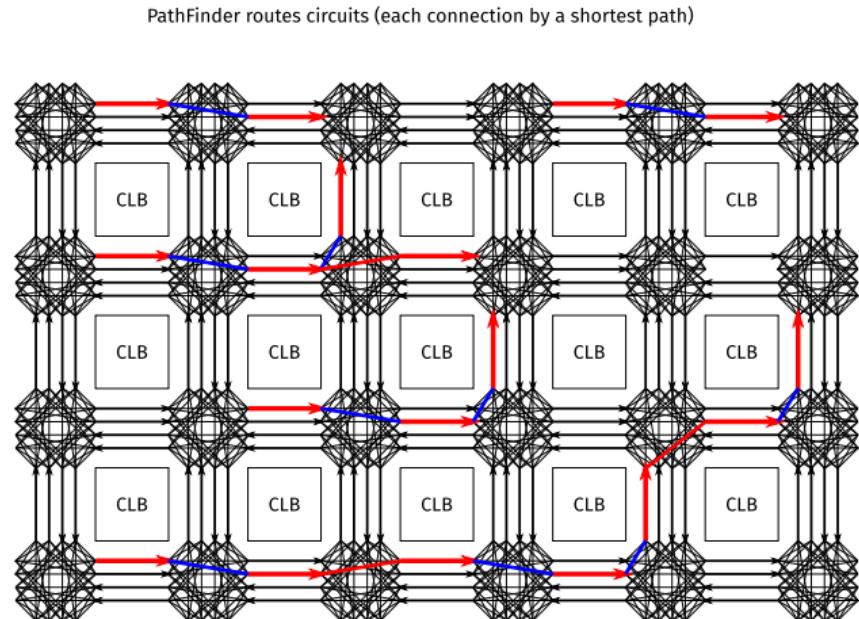
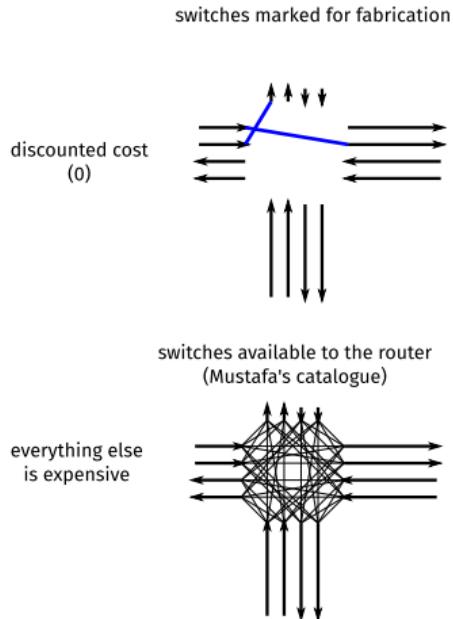
Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]



PathFinder routes circuits (each connection by a shortest path)



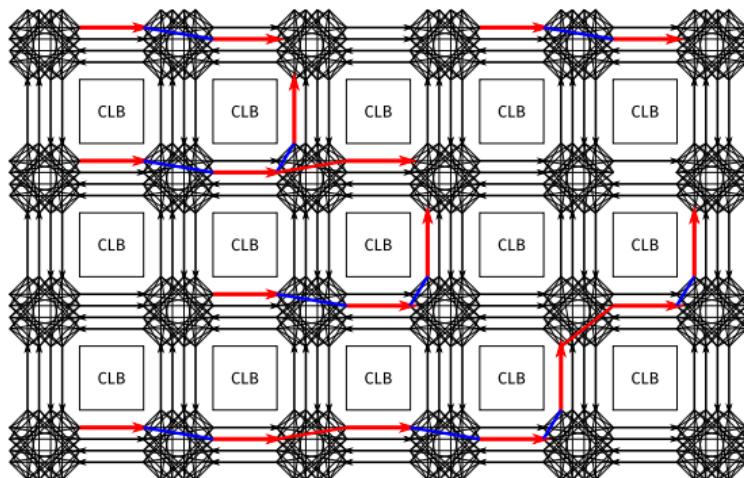
Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]



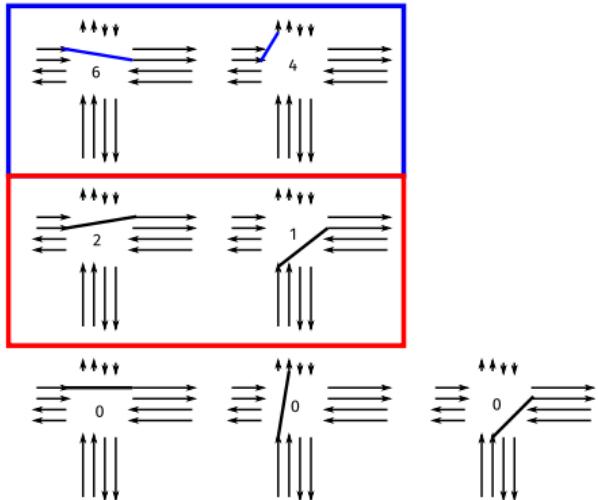
1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done

Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]

PathFinder routes circuits (each connection by a shortest path)



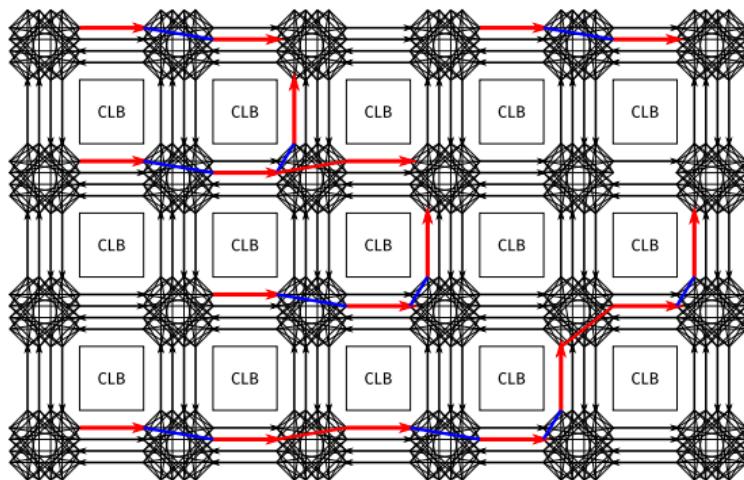
Count in how many switch-blocks each switch was used



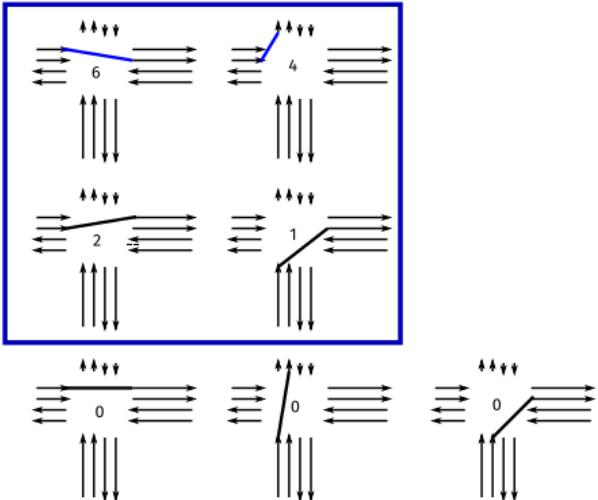
1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done
5. mark n most-used unmarked switches and set their cost to 0

Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]

PathFinder routes circuits (each connection by a shortest path)



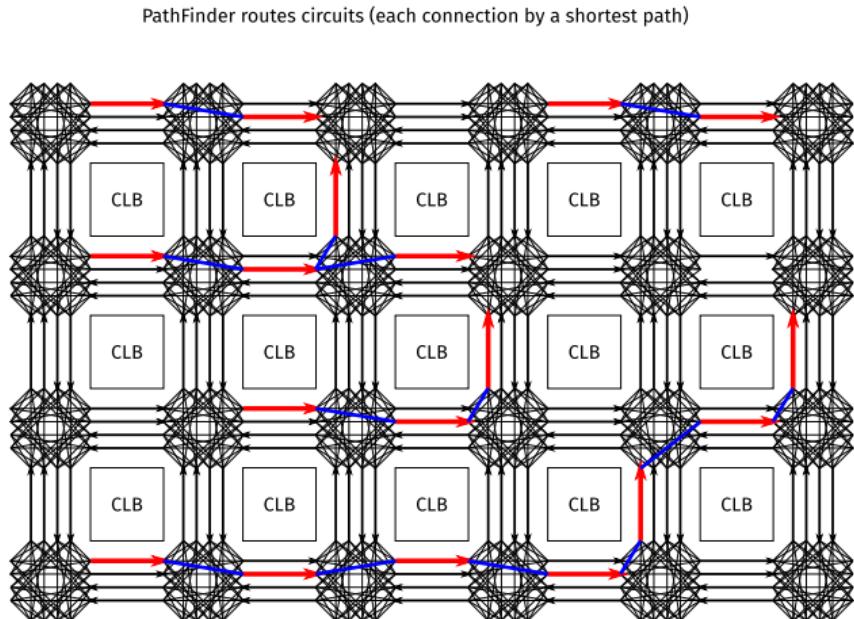
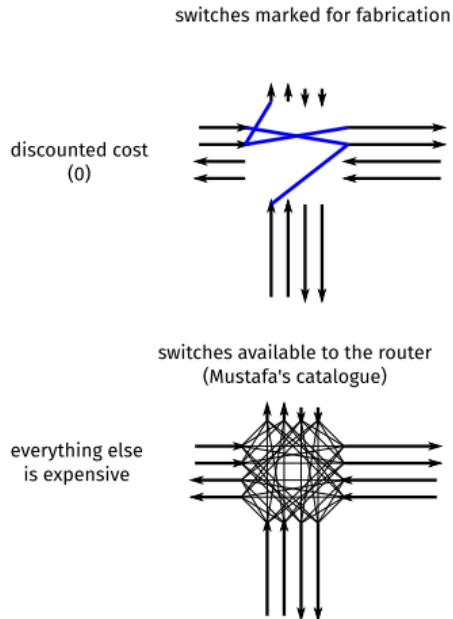
Give discounts and mark for fabrication



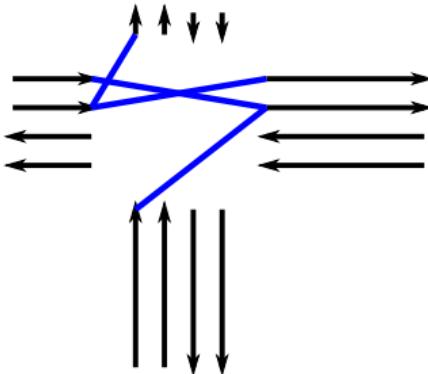
1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done
5. mark n most-used unmarked switches and set their cost to 0
6. goto 3

1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits

Avalanche Search [Nikolić and lenne. Turning PathFinder Upside-Down. FPL'21]



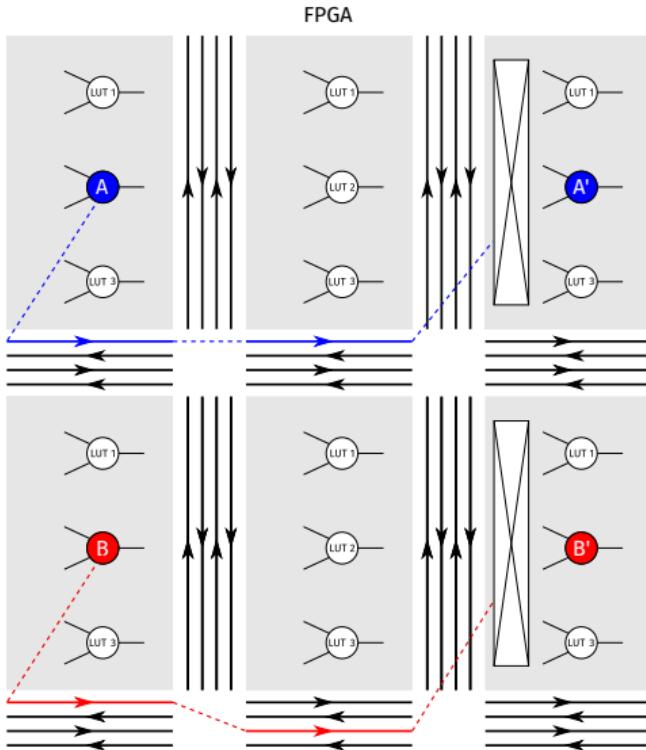
1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done



Final switch-pattern

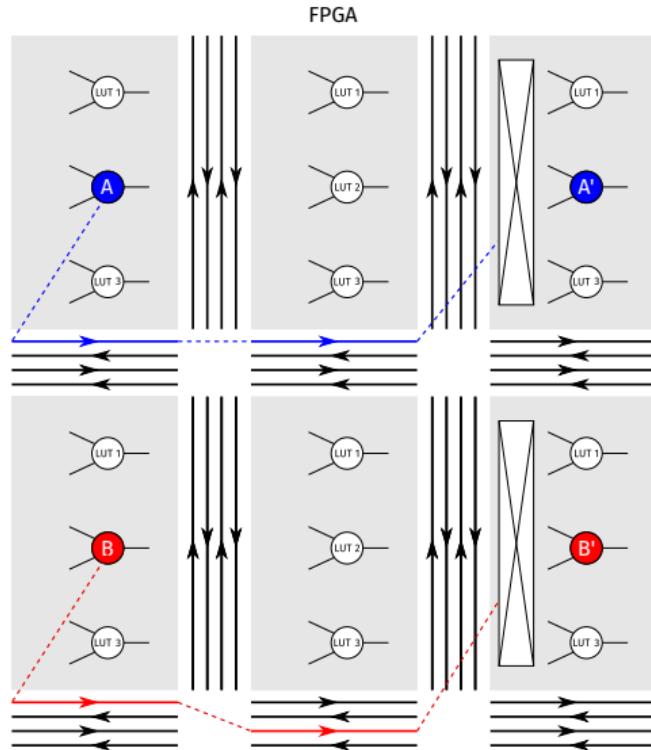
Where is the “avalanche”?

Avalanche costs



[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

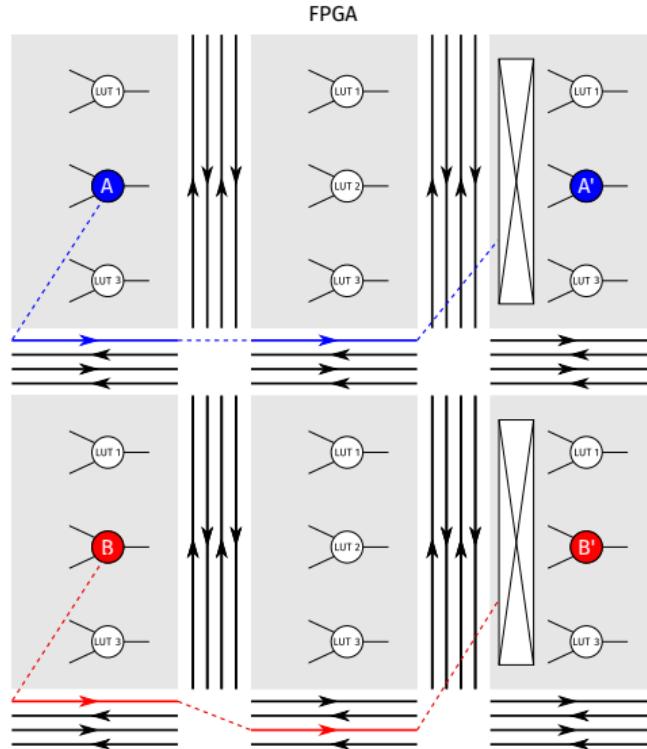
Avalanche costs



No congestion ✓

[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

Avalanche costs



No congestion ✓

But, switch-block must include both the straight and the slanted switch where one may suffice

How does PathFinder achieve signal spreading?

Congestion cost:

- all wire **INSTANCES** (individual rr-graph nodes) are initially **CHEAP**
- their cost **INCREASES** “in proportion with” conflict magnitude

[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL’21, Best Paper Award

How does PathFinder achieve signal spreading?

We need concentration instead of spreading

⇒ apply the same principle in reverse

[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

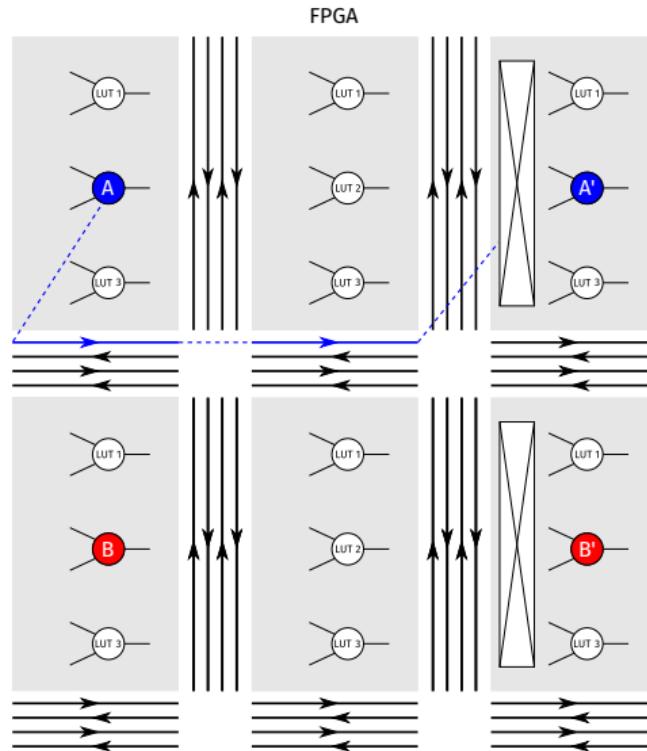
Avalanche costs

Avalanche cost:

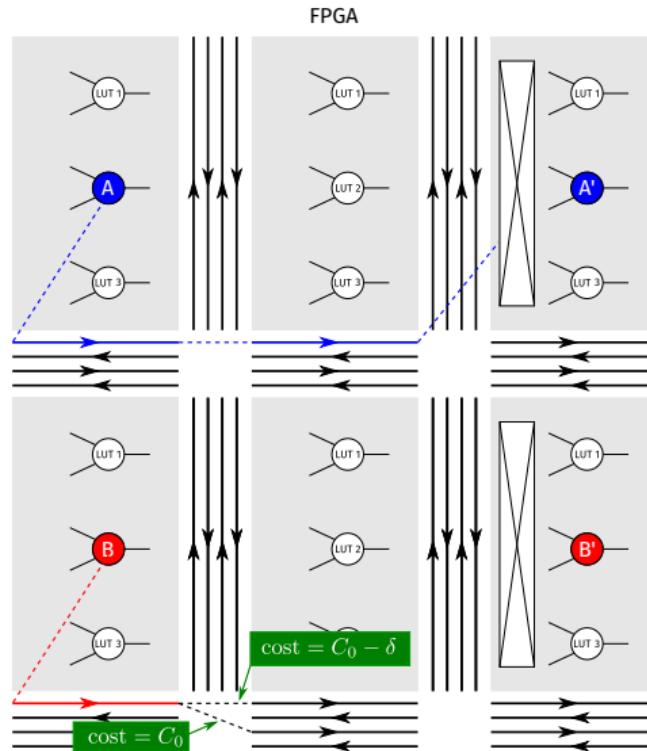
- all switches are initially **EXPENSIVE**
- their cost **DROPS** in proportion with how many of their instances are used

[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

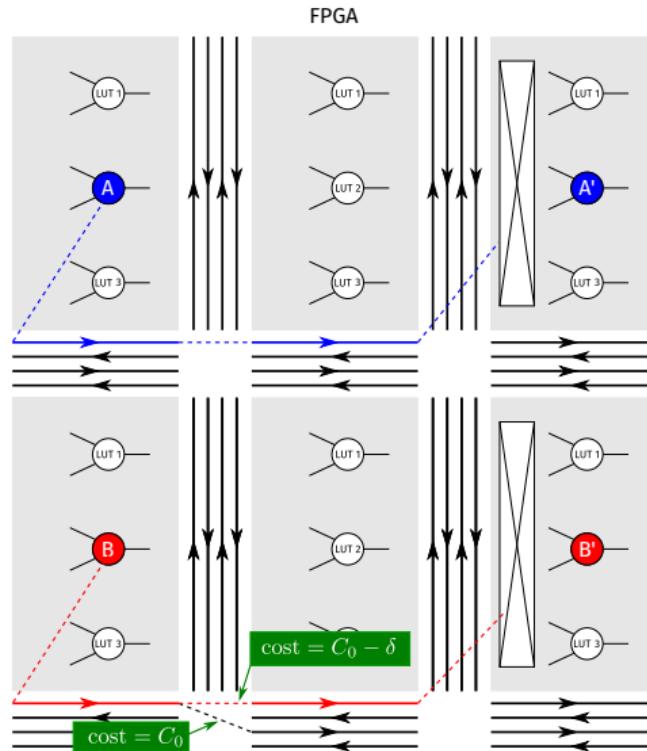
Avalanche costs



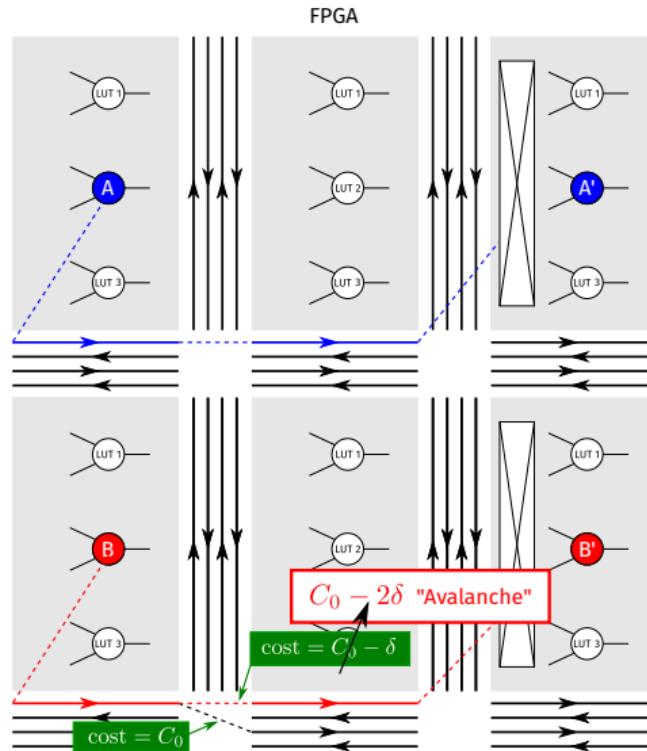
Avalanche costs



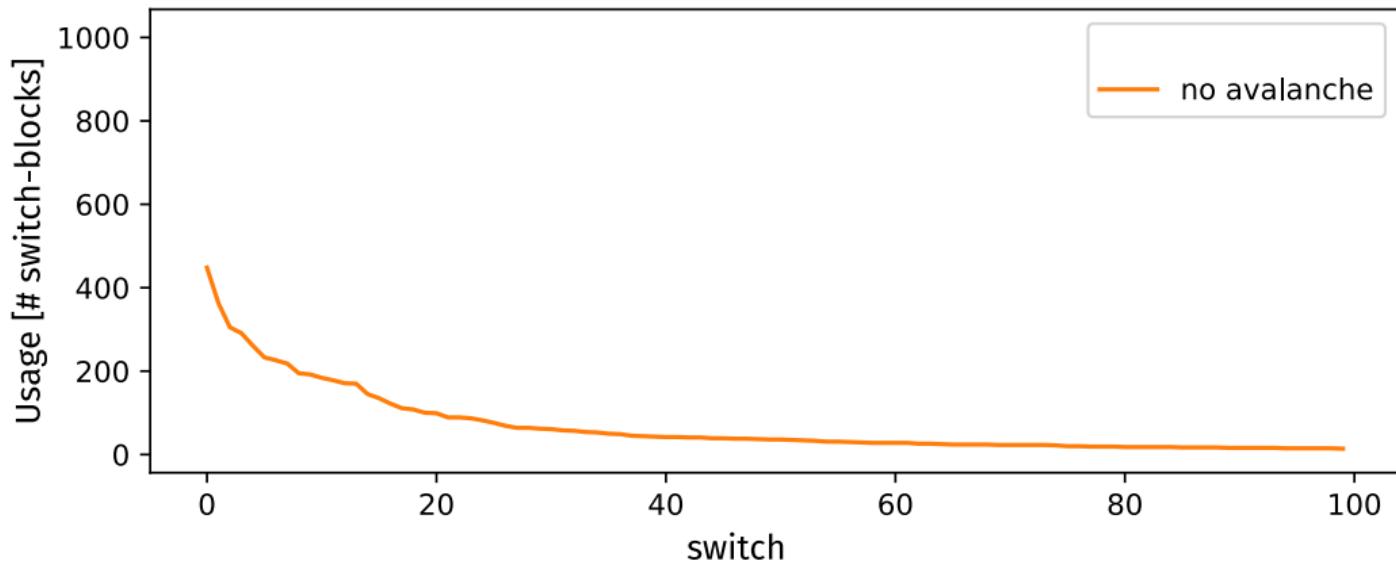
Avalanche costs



Avalanche costs

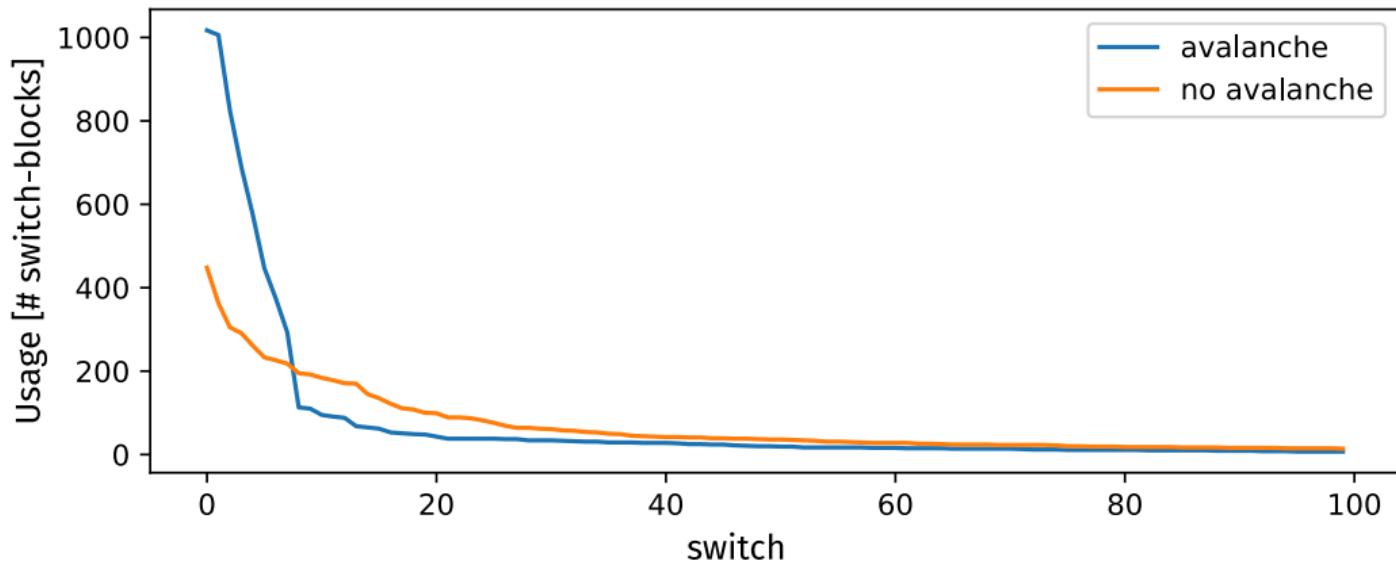


Avalanche costs: effect



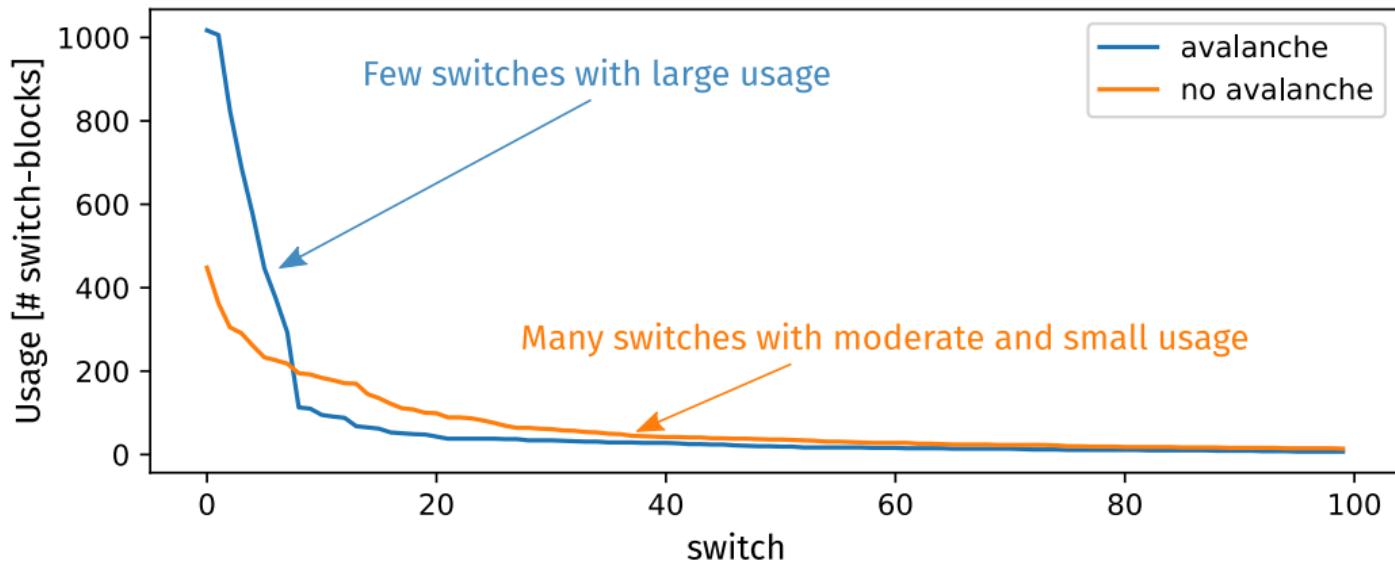
[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

Avalanche costs: effect



[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

Avalanche costs: effect



[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

Avalanche costs: function

Congestion cost: $b(u) \times (1 + p_{fac} \times O(u)) \times (1 + h_{fac} \times C_h(u))$

Avalanche cost: $\max(0, C(u) - (a_p \times U(u) + a_h \times U_h(u)))$

$U(u)$ —present number of SBs in which switch u is used

$U_h(u)$ —cumulative (historical) number of SBs in which switch u was used

[1] Nikolić and lenne. Turning PathFinder Upside Down, FPL'21, Best Paper Award

Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal

```
1: function CONGESTION_COST( $u, s$ )                                      ▷ computes the congestion cost of node  $u$  when routing signal  $s$ 
2:    if  $u \in RT(s)$  then return 0                                      ▷ if  $u$  is already used by one connection of  $s$ , it can freely be used by another
3:    return  $b(u) \times (1 + p_{fac} \times O(u)) \times (1 + h_{fac} \times C_h(u))$                                       ▷ otherwise, account for congestion
4: for  $u \in V$  do
5:     $O(u) = 0; C_h(u) = 0$                                       ▷ set occupancy and historical congestion of all nodes to 0
6:    if  $(\exists v \in V) ((u, v) \in E_c)$  then
7:        $RT(u) = \{u\}$                                               ▷ initialize the routing tree of each signal
8:     $i = 0; p_{fac} = p_{fac}^{init}$ 
9: do
10:   if  $i \geq \text{max\_iter}$  then return UNROUTABLE
11:   for  $s \in \{u \in V : (\exists v \in V) ((u, v) \in E_c)\}$  do
12:
13:     for  $u \in RT(s)$  do
14:        $O(u) = O(u) - 1$                                       ▷ reduce the occupancy of all nodes used by the signal  $s$  that is ripped up
15:        $RT(s) = \{s\}; O(s) = O(s) + 1$                               ▷ rip up the signal
16:       for  $t \in V : (s, t) \in E_c$  do
17:           $P = \text{SHORTEST\_PATH}(s, t, \forall u \in V : \text{cong}(u) = \text{CONGESTION\_COST}(u, s))$                       ▷ (re)route the connection  $s \rightarrow t$ 
18:          for  $u \in P$  do
19:           if  $\neg(u \in RT(s))$  then
20:               $O(u) = O(u) + 1$                                       ▷ increase the occupancy of all nodes not already used by the signal  $s$ 
21:               $RT(s) = RT(s) \cup P$                                       ▷ add the connection route to the routing tree of  $s$ 
22:       for  $u \in V$  do
23:         $C_h(u) = C_h(u) + \max(0, O(u) - 1)$                               ▷ update historical congestion
24:         $i = i + 1$ 
25:        $p_{fac} = p_{fac} \times p_{fac}^{mult}$                               ▷ increase the penalty of using occupied nodes;  $p_{fac}^{mult} > 1$  (1.3 is default in VPR [16])
26:       while  $\exists u \in V : O(u) > 1$                                       ▷ finish if there is no congestion
27:       return  $\forall RT$                                                       ▷ return all routing trees
```

Update $U(u)$ as well

iterations
in iteration;
in this [16]

Algorithm 1 Simplified PathFinder [McMurchie95, Betz98]

Input: $G = (V, E)$ —rr-graph, $E_c \subseteq V \times V$ —all connections to be routed; **Output:** A routing tree of each signal s

```

1: function CONGESTION_COST( $u, s$ )           ▷ computes the congestion cost of node u when routing signal s
2:   if  $u \in RT(s)$  then return 0          ▷ if u is already used by one connection of s, it can freely be used by another
3:   return  $b(u) \times (1 + p_{fac} \times O(u)) \times (1 + h_{fac} \times C_h(u))$           ▷ otherwise, account for congestion

4: for  $u \in V$  do
5:    $O(u) = 0; C_h(u) = 0$                   ▷ set occupancy and historical congestion of all nodes to 0
6:   if  $(\exists v \in V) ((u, v) \in E_c)$  then
7:      $RT(u) = \{u\}$                       ▷ initialize the routing tree of each signal
8:    $i = 0; p_{fac} = p_{fac}^{init}$ 
9:   do
10:    if  $i \geq \text{max\_iter}$  then return UNROUTABLE  ▷ no congestion-free routing was found in max_iter iterations
11:    for  $s \in \{u \in V : (\exists v \in V) ((u, v) \in E_c)\}$  do  ▷ all signals are ripped up and rerouted in each iteration;
12:      for  $u \in RT(s)$  do  ▷ modern incremental routers deviate from this [16]
13:         $O(u) = O(u) - 1$                   ▷ reduce the occupancy of all nodes used by the signal s that is ripped up
14:         $RT(s) = \{s\}; O(s) = O(s) + 1$           ▷ rip up the signal
15:        for  $t \in V : (s, t) \in E_c$  do
16:           $P = \text{SHORTEST\_PATH}(s, t, \forall u \in V : \text{cong}(u) = \text{CONGESTION\_COST}(u, s))$  ▷ (re)route the connection  $s \rightarrow t$ 
17:          for  $u \in P$  do
18:            if  $\neg(u \in RT(s))$  then
19:               $O(u) = O(u) + 1$ 
20:             $RT(s) = RT(s) \cup P$           ▷ add the connection route to the routing tree of s
21:          for  $u \in V$  do
22:             $C_h(u) = C_h(u) + \max(0, O(u) - 1)$   ▷ update historical congestion
23:           $i = i + 1$ 
24:         $p_{fac} = p_{fac} \times p_{fac}^{mult}$   ▷ increase the penalty of using occupied nodes;  $p_{fac}^{mult} > 1$  (1.3 is default in VPR [16])
25:        while  $\exists u \in V : O(u) > 1$   ▷ finish if there is no congestion
26:      return VRT  ▷ return all routing trees

```

Update $U_h(u)$ as well

by used by the signal s

▷ add the connection route to the routing tree of s

▷ update historical congestion

▷ increase the penalty of using occupied nodes; $p_{fac}^{mult} > 1$ (1.3 is default in VPR [16])

▷ finish if there is no congestion

▷ return all routing trees

Update $U_h(u)$ as well

$\forall u \in V$ do

$$C_h(u) = C_h(u) + \max(0, O(u) - 1)$$

- ▶ add the connection route to the routing tree of s

► update historical congestion

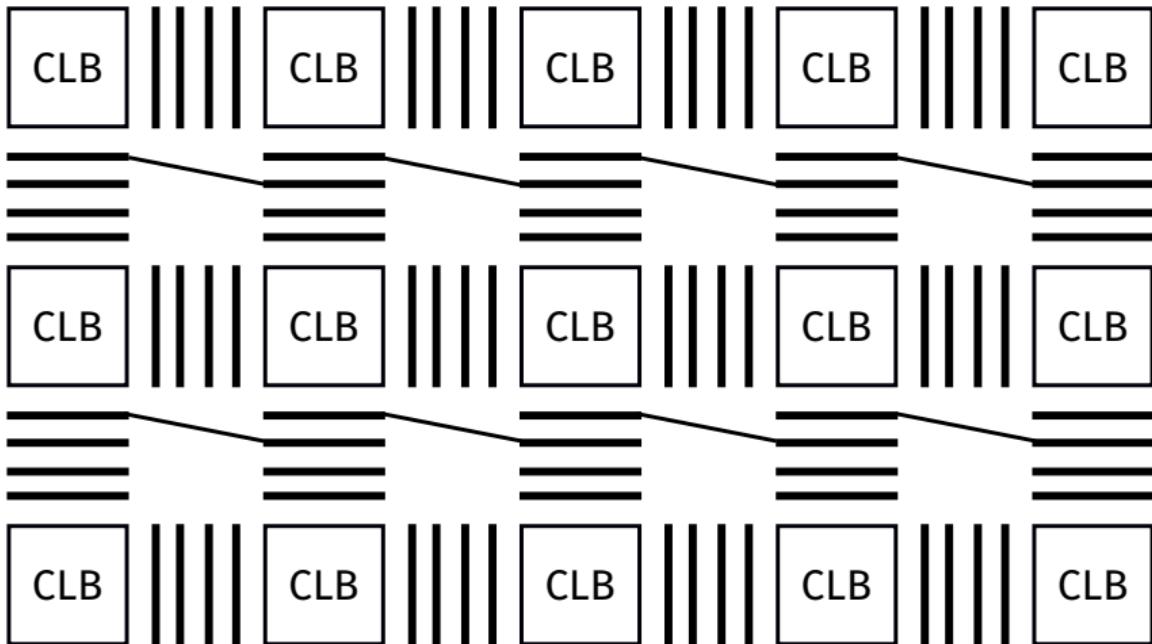
Avalanche vs congestion costs

Congestion

- **UNIQUE** for **EACH** wire instance
- **PENALIZES** using same **WIRE INSTANCES** by different nets

Avalanche

- **COMMON TO ALL** instances of the same switch
- **AWARDS** using same **SWITCH TYPES** by different nets



All instances of the same switch share the same avalanche cost

Avalanche vs congestion costs

Congestion

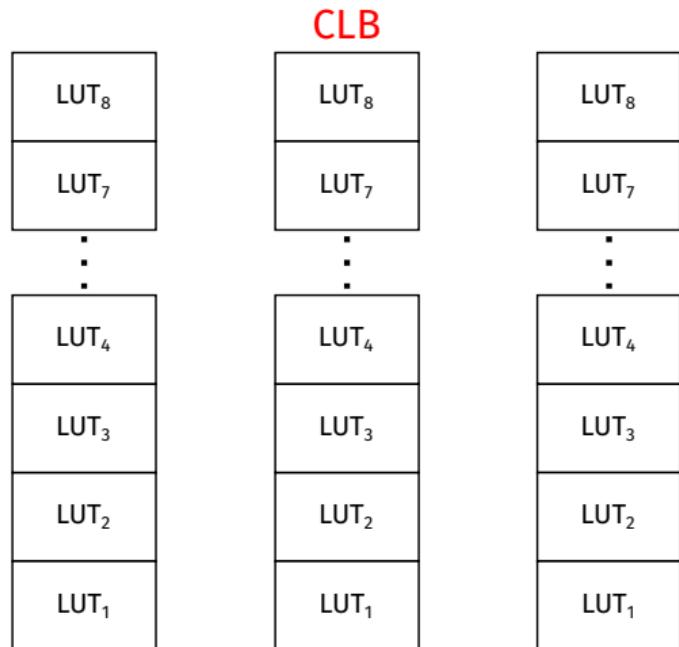
- **UNIQUE** for **EACH** wire instance
- **PENALIZES** using same **WIRE INSTANCES** by different nets

Avalanche

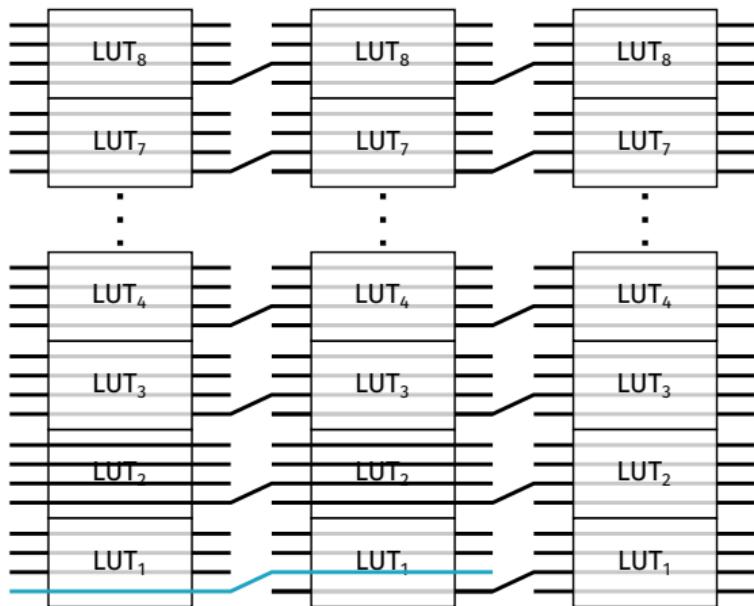
- **COMMON TO ALL** instances of the same switch
- **AWARDS** using instances of the same **SWITCH** by different nets

Effectiveness of Avalanche Search

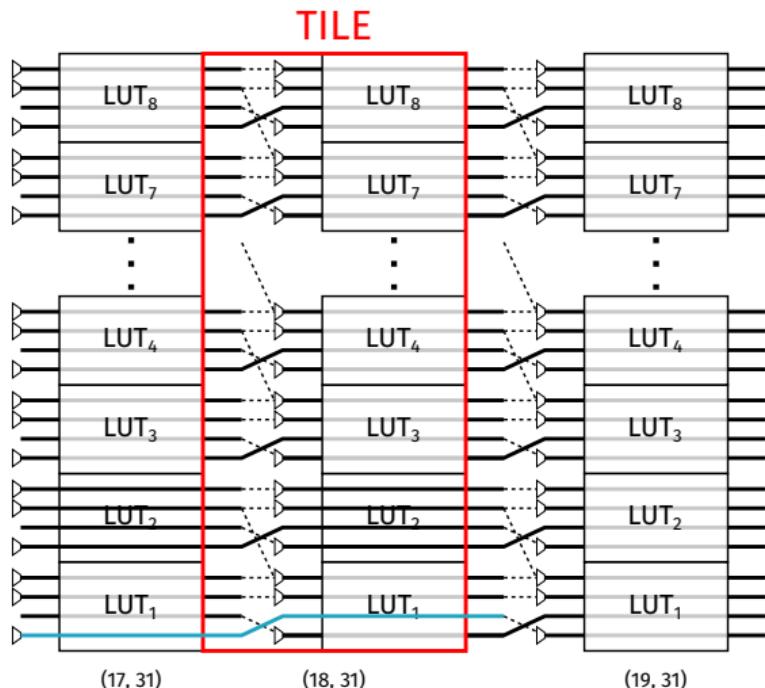
A more physical view of Island-Style FPGAs: stacked LUTs



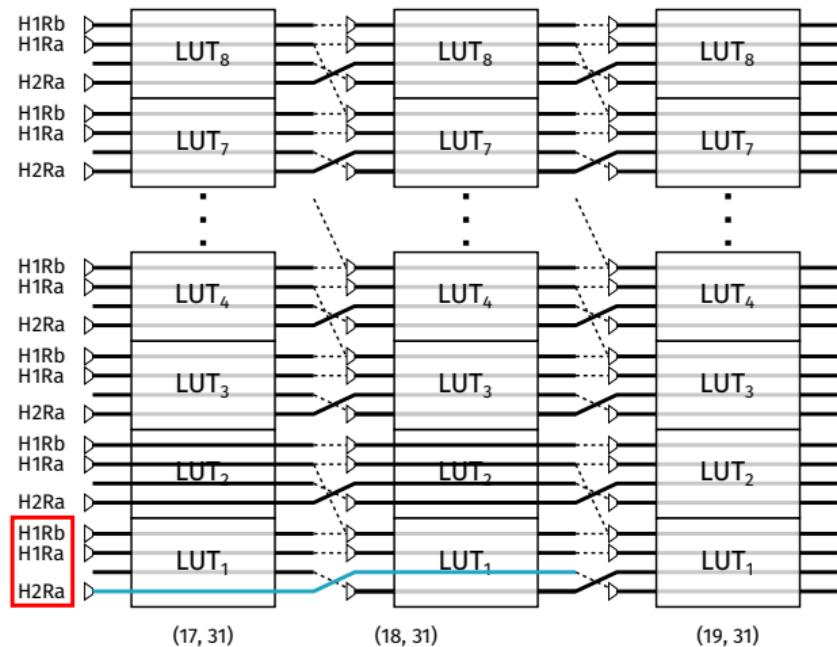
A more physical view of Island-Style FPGAs: wires on top



A more physical view of Island-Style FPGAs: muxes on the side

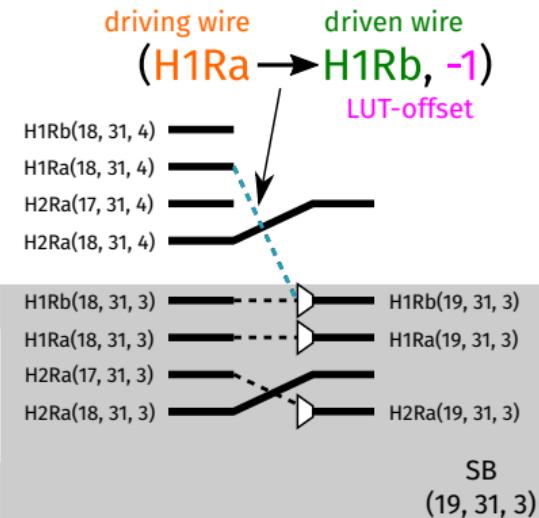
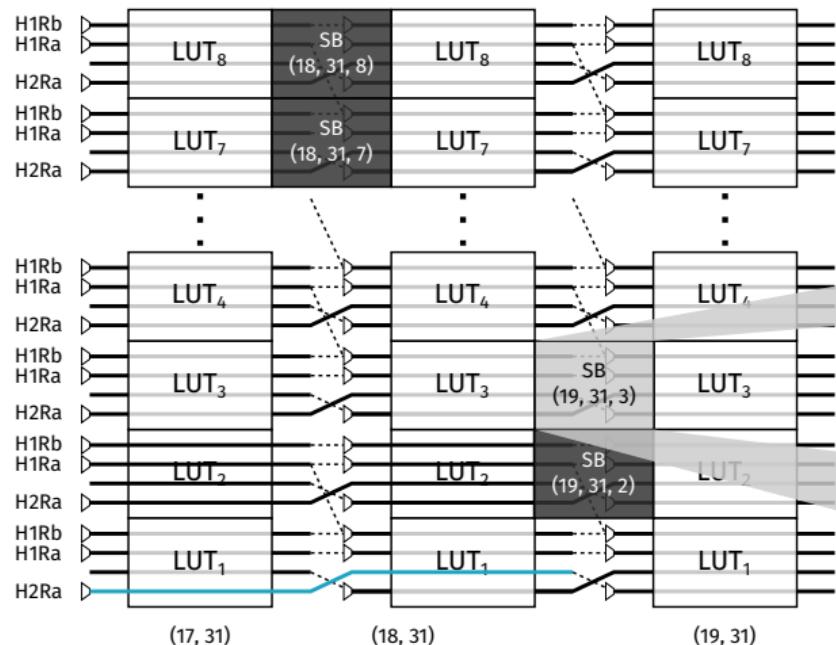


A more physical view of Island-Style FPGAs: distributed channels



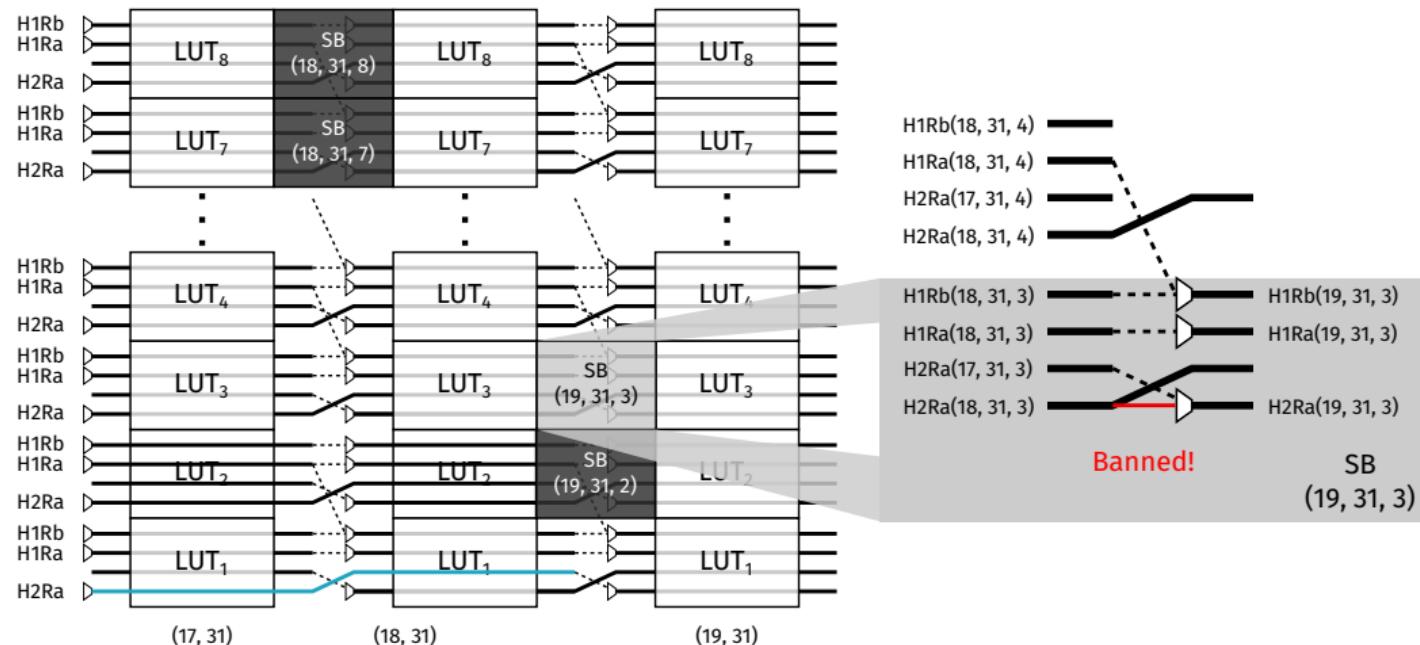
[1] Chromczak et al. Architectural Enhancements in Intel Agilex FPGAs. FPGA'20

A more physical view of Island-Style FPGAs: distributed SBs



[1] Chromczak et al. Architectural Enhancements in Intel Agilex FPGAs. FPGA'20
(our guess)

A more physical view of Island-Style FPGAs: no intermediate taps



Our best manual effort (4nm; 16 wires; 564 possible switches)

$\langle F_S \rangle = 11.25$

H1La - 0	3	1	1	1										
H1Lb -	3	0	1	1	1									
H2La -	1	1	1	1	1									
H4La -	1	1	1	1	1									
H6La -	1	1	1	1	1									
H1Ra -		0	3	1	1	1	0	1	1	0	1	1		
H1Rb -		3	0	1	1	1	1	0	1	1	0	1		
H2Ra -		1	1	1	1	1	1	1	1	1	1	1		
H4Ra -		1	1	1	1	1	1	1	1	1	1	1		
H6Ra -		1	1	1	1	1	1	1	1	1	1	1		
V1Ua - 0	0	1	1	1	1	0	1	1	1	0	3	1		
V1Ub -	1	0	1	1	1	1	0	1	1	1	3	0	1	
V4Ua -	1	1	1	1	1	1	1	1	1	1	1	1		
V1Da - 0	0	1	1	1	1	0	1	1	1	1	0	3	1	
V1Db -	1	0	1	1	1	1	0	1	1	1	3	0	1	
V4Da -	1	1	1	1	1	1	1	1	1	1	1	1	1	



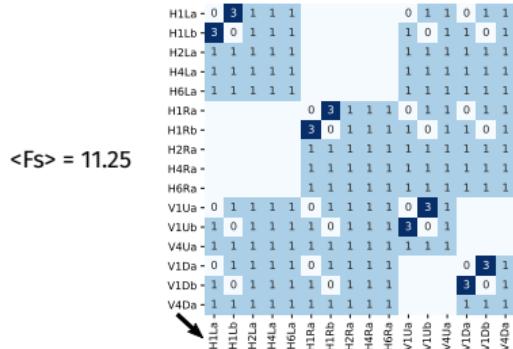
#switches	180		
	fi	fo	t[ps]
average →			
H1	10	10	16.0
H2	11	11	21.3
H4	11	11	30.8
H6	11	11	43.1
V1	12	12	24.6
V4	13	13	74.3
W(tile)	7464 nm		
CPD	1.46 ns		

[1] Nikolić, Catthoor, Tőkei, and lenne. Global Is the New Local. FPGA'21

Avalanche Search

- 3 MCNC circuits used in exploration
- 17 MCNC circuits used for testing
- converged in 36 iterations

Avalanche Search



#switches	180		
average →	fi	fo	t[ps]
H1	10	10	16.0
H2	11	11	21.3
H4	11	11	30.8
H6	11	11	43.1
V1	12	12	24.6
V4	13	13	74.3
W(tile)	7464 nm		
CPD	1.46 ns		

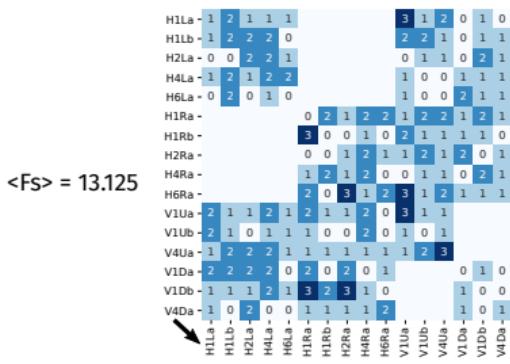
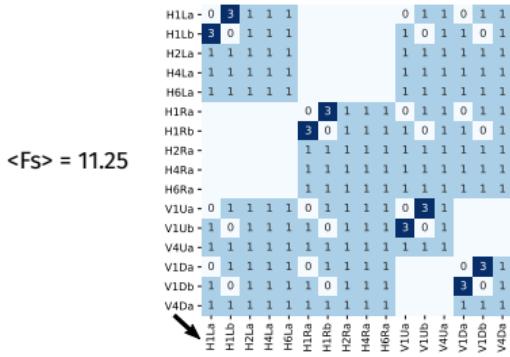
#switches	78		
average →	fi	fo	t[ps]
H1	5	3	13.9
H2	5	4	16.8
H4	4	7	27.4
H6	5	5	35.7
V1	7	6	21.8
V4	2	5	70.1
W(tile)	6792 nm -9%		
CPD	1.38 ns -5.48%		

Simulated annealing

- inspired by [1]
- each move is addition or removal of a single switch
- cost function is a combination of tile area and
geomean routed critical path delay of the same 3 MCNC circuits
- low-temperature anneal of the manual switch-pattern
- 100 temperature changes
- 100 moves per temperature

[1] Lin, Wawrynek, El Gamal. Exploring FPGA routing architecture stochastically. TCAD'10

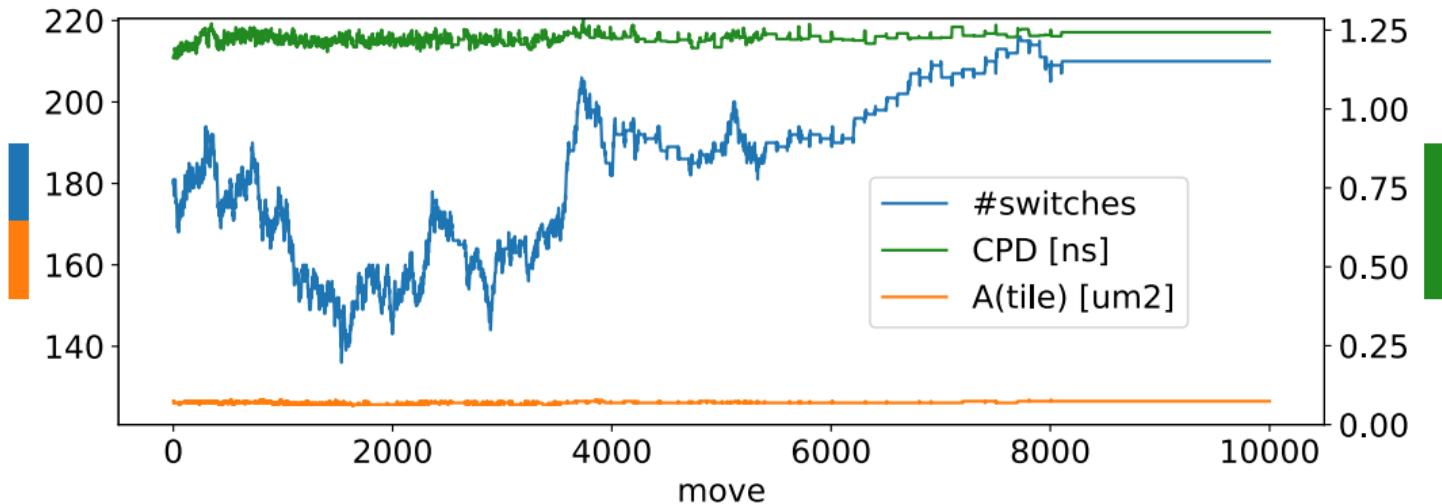
Simulated annealing



#switches	180		
	average →	fi	fo
	t[ps]		
H1	10	10	16.0
H2	11	11	21.3
H4	11	11	30.8
H6	11	11	43.1
V1	12	12	24.6
V4	13	13	74.3
W(tile)	7464 nm		
CPD	1.46 ns		

#switches	210		
	average →	fi	fo
	t[ps]		
H1	13	13	19.6
H2	14	11	24.1
H4	16	12	32.1
H6	9	13	47.3
V1	14	15	29.2
V4	13	15	86.8
W(tile)	7488 nm +0.32%		
CPD	1.55 ns +6.16%		

Simulated annealing



With a better cost function, directed moves, etc., it would likely work better

Routability

Testing on

- 10 Gnl [1] circuits
- each with $\sim 10\,000$ LUTs
- all with Rent's exponent 0.7

[3] Stroobandt, Depreitere, and Campenhout. Generating new benchmark designs.
Integration'99

Routability

Congested nodes after 300 PathFinder iterations [%]

circuit	1	2	3	4	5	6	7	8	9	10
avalanche	0.906%	0.274%	0%	0.521%	0.323%	0.187%	0.149%	0.262%	0.040%	0.002%

- only 1/10 circuits was routed
- sparsification of the switch-pattern was enabled by MCNC circuits not requiring the channel width dimensioned for larger circuits

Routability: remedy

Perform additional iterations of Avalanche Search,
this time using one Gnl circuit per iteration

Routability: remedy

H1La -	1	0	0	0	0	0	0	0	1	0	0
H1Lb -	1	0	1	0	0	0	0	1	1	0	0
H2La -	0	0	2	0	0	0	1	0	1	1	0
H4La -	1	1	1	1	1	1	1	1	0	0	0
H6La -	0	1	0	0	1	1	1	0	0	0	1
H1Ra -		0	1	0	0	0	0	1	0	0	0
H1Rb -		1	0	0	0	1	0	1	1	1	1
H2Ra -		0	1	0	0	0	1	1	0	1	1
H4Ra -		1	1	1	0	1	1	1	0	0	1
H6Ra -		0	2	1	0	1	1	0	0	1	0
V1Ua -	1	0	1	1	1	0	1	1	1	1	1
V1Ub -	0	1	0	0	0	0	1	0	0	1	1
V4Ua -	0	0	0	1	0	0	1	0	0	0	1
V1Da -	0	0	1	0	0	0	1	1	0	1	0
V1Db -	0	1	0	0	1	1	1	0	0	1	0
V4Da -	1	0	1	1	0	0	1	1	0	0	1

Number of iterations until congestion is resolved

circuit	1	2	3	4	5	6	7	8	9	10
Gnl-extended avalanche	142	61	27	106	26	55	46	55	30	82

- converged in 3 iterations
- added 6 more switches (now 84)
- routed delay went up from 1.38 to 1.39 ns
- all 10 Gnl circuits are routable

Regularizing Avalanche Search

Layout prefers regularity



United States Patent [19]

[11] Patent Number: 5,744,995

[14] Date of Patent: Apr. 28, 1998

[54] SIX-INPUT MULTIPLEXER WITH TWO GATE LEVELS AND THREE MEMORY CELLS

5,841,480 5/1998 Shieh et al. 325/404
5,845,574 7/1998 Nomura 325/407
5,842,530 7/1998 Chang et al. 325/408
5,870,851 10/1998 Cheng et al. 325/407

[73] Inventor: Steven P. Young, San Jose, Calif.

Primary Examiner—Timothy P. Callahan
Assistant Examiner—My-Trang Nhu Ton

[73] Assignee: Xilinx, Inc., San Jose, Calif.

Attorney, Agent, or Firm—Eliel M. Young; Adam H. Tschirhart

[21] Appl. No.: 658,896

[22] Filed: Apr. 17, 1996

[15] Int. Cl. 4 H03K 17/00

[51] U.S. Cl. 325/407, 325/408

[58] Field of Search 327/99, 407, 416,

327/410, 411, 413, 427, 437, 326/38, 39,

40, 41, 44, 45, 49

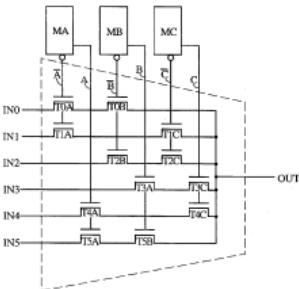
[56] References Cited
U.S. PATENT DOCUMENTS

3,614,327 10/1971 Low 325/405
4,535,344 8/1985 Chang et al. 325/407
4,932,634 6/1990 Peng et al. 325/408
5,050,841 7/1991 Hellmann et al. 325/409

ABSTRACT

A six-input multiplexer is disclosed using only two transistors in the signal path from an input port to the output port. The multiplexer has three control inputs and two output ports. The multiplexer uses three control signals and requires that the control signal combinations 000 and 111 not be used. The control signals are A, B, and C. The control signals 000, 100, 101, and 110 can be used to select between six input signals by placing only two transistors in the signal path. The control signals 111 and 011 are not used because the output signals are the same and the third bit is different from the other two. A compact layout results when two multiplexers use common input signals.

6 Claims, 7 Drawing Sheets



U.S. Patent

Apr. 28, 1998

Sheet 6 of 7

5,744,995

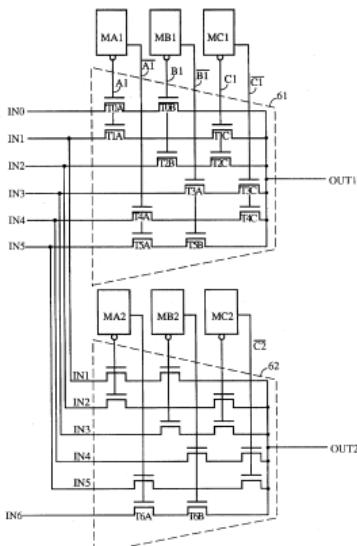
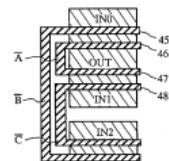
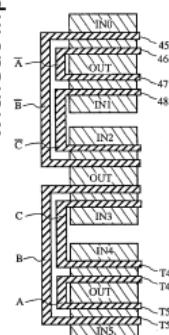


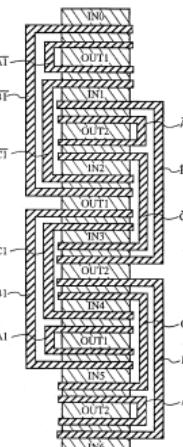
FIG. 7



without input sharing



with input sharing



Layout prefers regularity

If multiplexers are 6:1:

- fewer memory cells are needed (normal 2-level 6:1 needs 4 SRAMs)

If two multiplexers share 5/6 inputs:

- much more compact layout (diffusion sharing)
- fewer vias
- shorter wiring

Layout prefers regularity

Can we constrain Avalanche Search to produce only regular solutions?

1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if no unmarked switches are used, done
5. mark n most-used unmarked switches and set their cost to 0
6. goto 3

Key points: constructing a feasible solution

After step 3 (PathFinder routing):

1. Encode **ANY** regularity constraints
(e.g., for layout or CAD tools)
as an Integer Linear Program (ILP)
2. Let ILP maximize PathFinder's "desire" (usage of different switches)
while satisfying the above constraints

Key points: updating costs

For switches in the ILP solution (satisfying regularity constraints)

1. Mark n most-used unmarked switches and reduce their cost to 0

For switches not in the ILP solution (violating regularity constraints)

1. Assign full cost with no discount

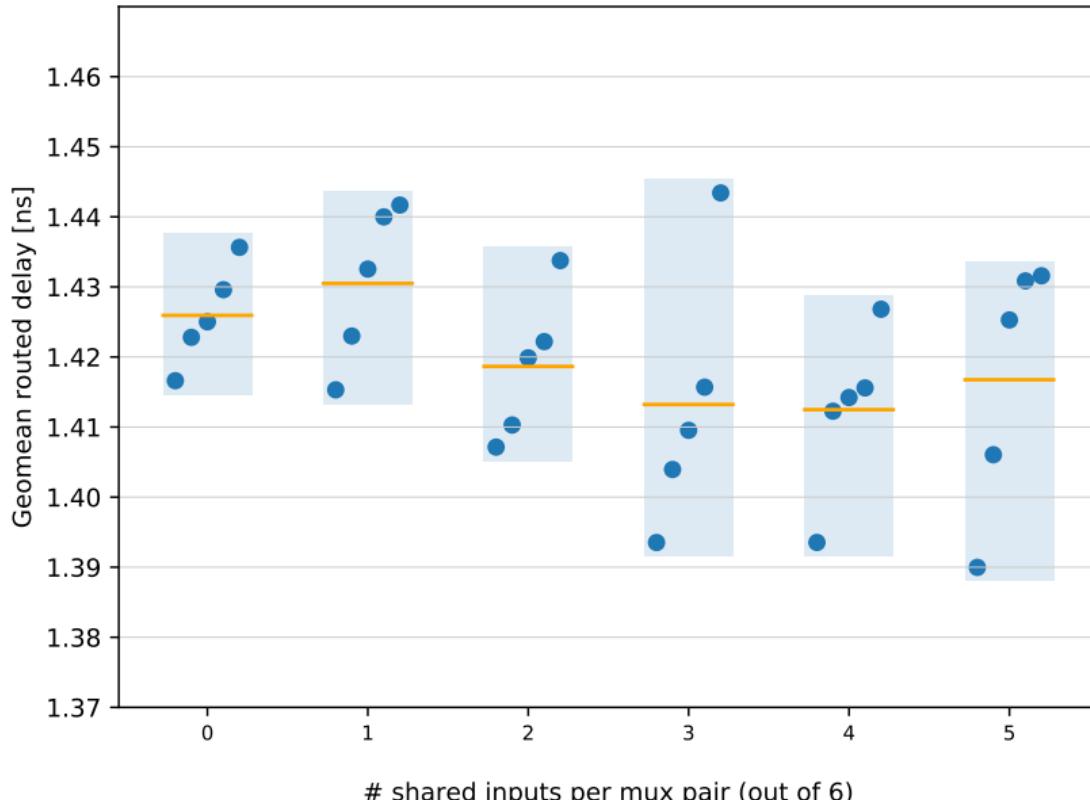
Key points: ensuring that constraints are met

- Final pattern is the last ILP solution

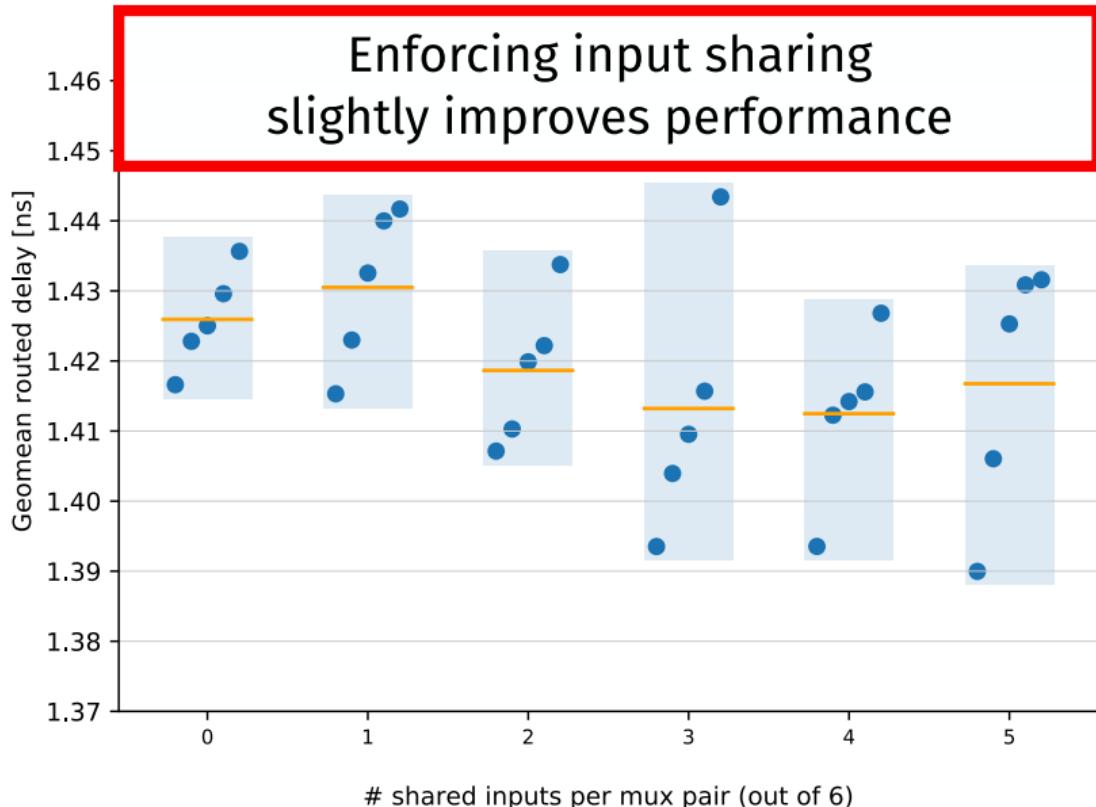
Regularizing Avalanche Search

1. switches marked for fabrication $\leftarrow \{\}$
2. all possible switches can be used at cost C
3. let PathFinder route the circuits
4. if iterations expanded \implies return the last ILP solution
5. solve the ILP, always retaining marked switches
6. mark n most-used unmarked switches from ILP's solution
and set their cost to 0
7. goto 3

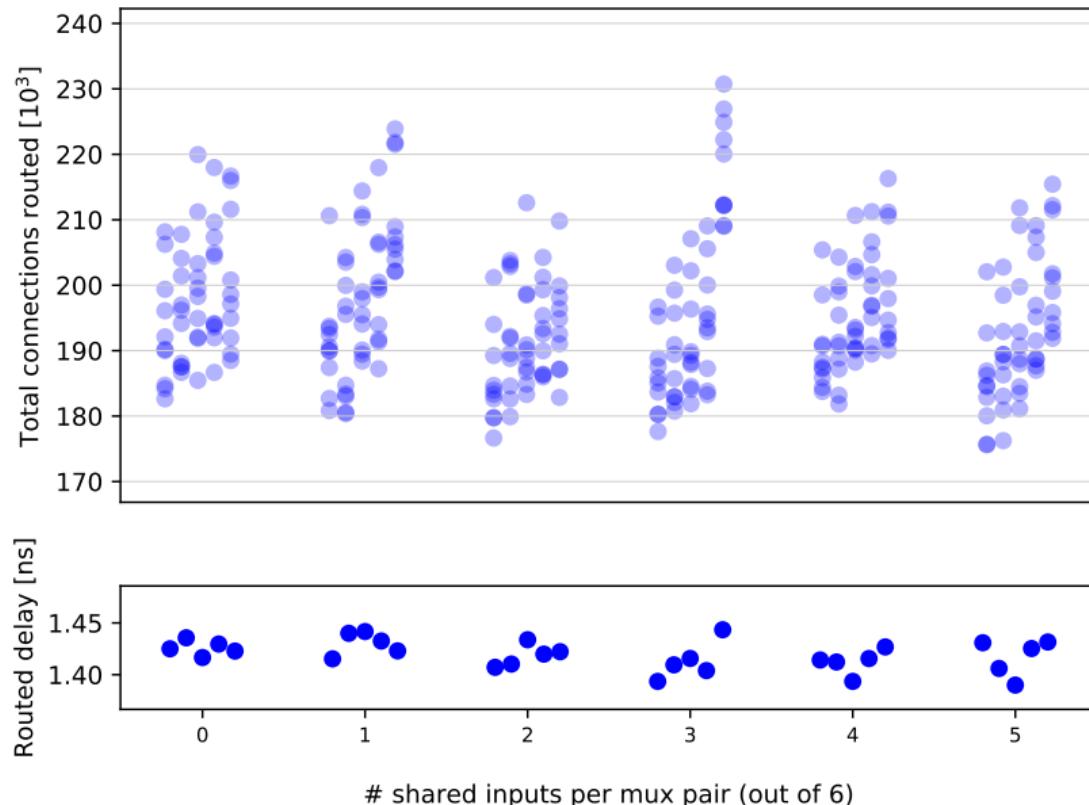
Input sharing: delay



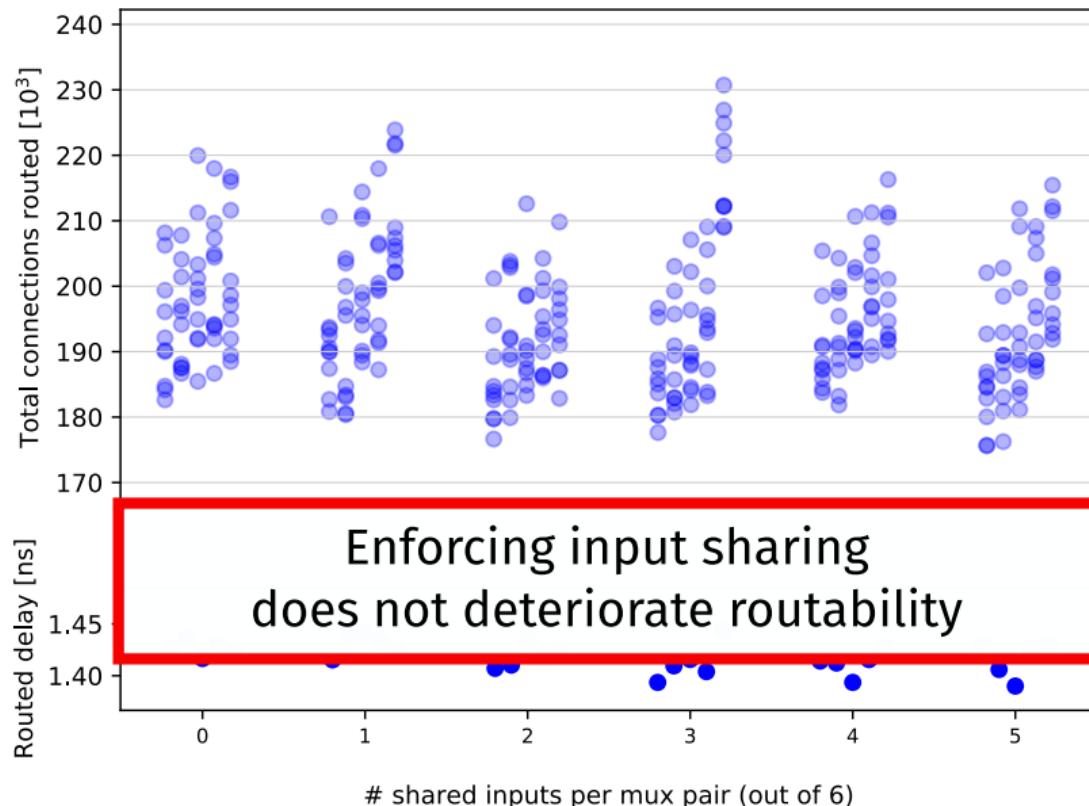
Input sharing: delay



Input sharing: routability



Input sharing: routability



Runtime scalability

Routing graph size

- embedding the entire design space in the routing graph increases its size significantly
- this slows down shortest path search

- avalanche costs can drop to zero
- for the lookahead to be admissible, we must compute it for all avalanche costs at zero
- this makes it very ineffective when switch usage is low (costs $\sim 1000 \times$ base costs)

A*: possible remedy

- store the minimum number of unmarked nodes to reach each tile
- multiply this number by the current lowest cost of an unmarked switch
- maybe some more preprocessing can be afforded as avalanche costs are common to all instances of the same switch, unlike congestion

Conclusions

- we now have a method for automatically designing switch-patterns without explicitly listing and testing solutions (previously a fundamental bottleneck)
- we can also make the solutions respect arbitrary constraints
- because the method operates on a routing graph, it could be useful for other aspects of interconnect (e.g., intracluster connectivity, channel segmentation, etc.)
- more work is needed to resolve the A^* issues